Microsoft

LearnAI - Machine Learning on Azure

Seth Mottaghinejad
Wolfgang Pauli, PhD

LearnAI@Microsoft.com

Overview

- Machine Learning on Azure
- Custom AI
- Compute Targets (DSVMs and Managed Compute)
- DevOps for Machine Learning
- Azure Machine Learning Pipelines
- Flexible and Support for Open Source Frameworks
- Deployment
- Tool Agnostic Python SDK

Building blocks for a Data Science Project

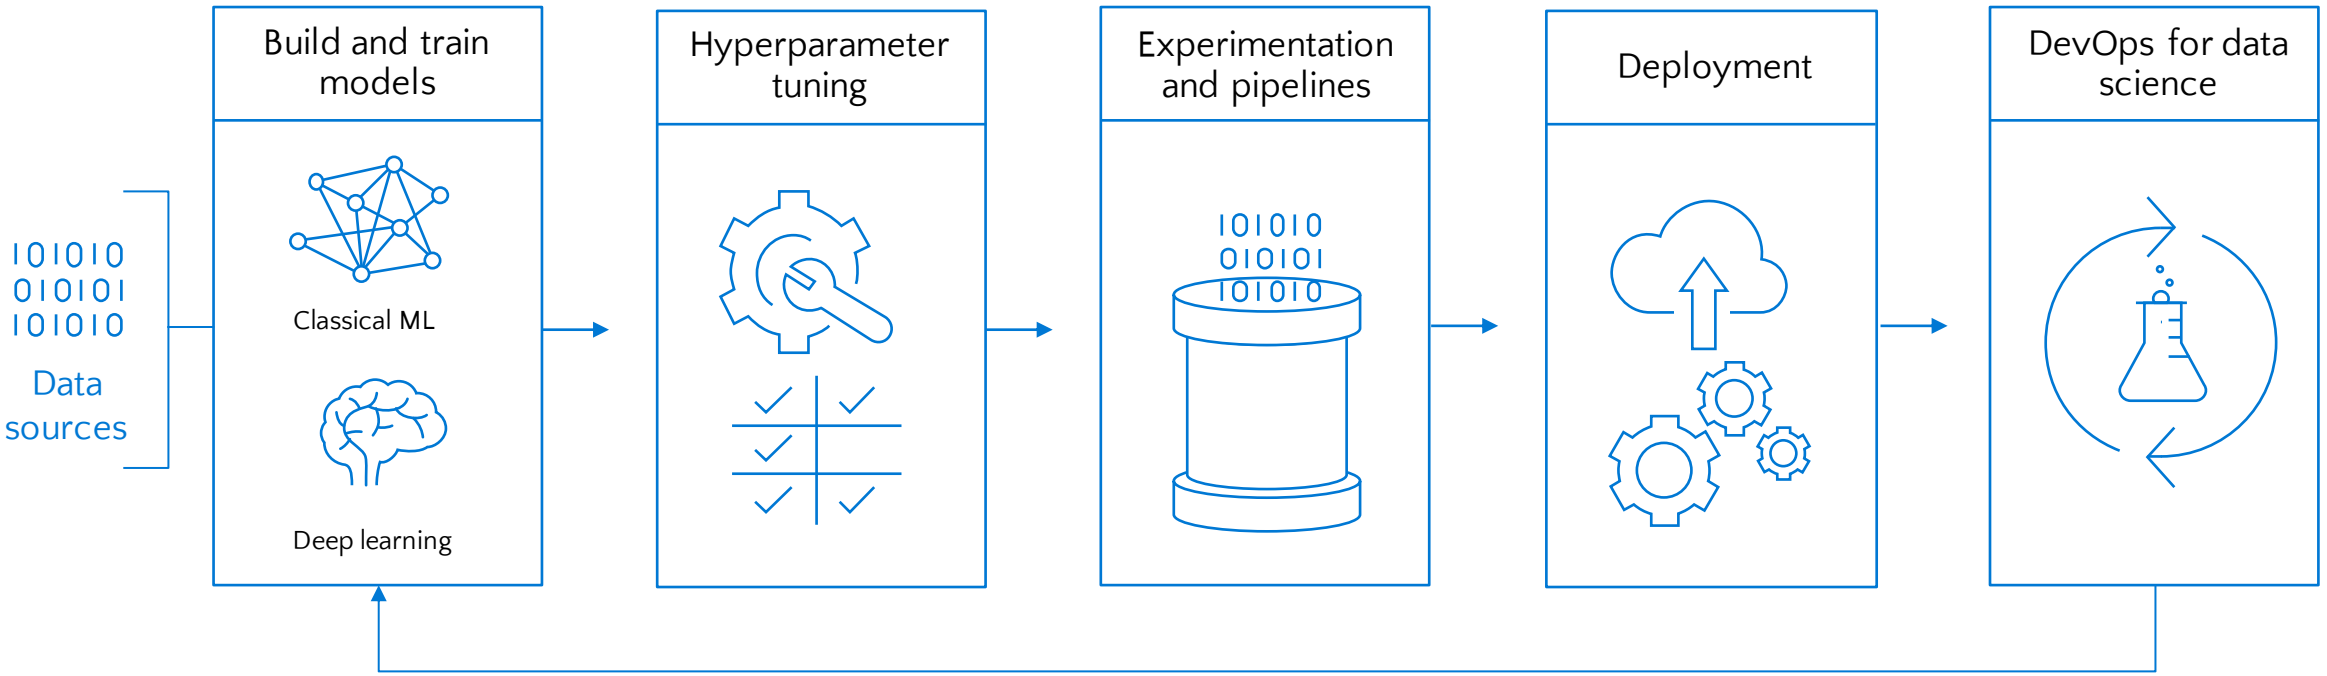

Machine Learning on Azure

Domain Specific Pretrained Models
To reduce time to market

- Vision
- Speech
- Language
- Search

Familiar Data Science Tools
To simplify model development

- PyCharm
- Jupyter
- Visual Studio Code
- Command line

Popular Frameworks
To build machine learning and deep learning solutions

- PyTorch
- TensorFlow
- Scikit-Learn
- ONNX

Productive Services
To empower data science and development teams

- Azure Databricks
- Azure Machine Learning
- Machine Learning VMs

Powerful Hardware
To accelerate deep learning

- CPU
- GPU
- FPGA

From the Intelligent Cloud to the Intelligent Edge

Azure Machine Learning Service

Set of Azure Cloud Services

Python SDK

That enables you to:

- ✓ Prepare Data
- ✓ Build Models
- ✓ Train Models
- ✓ Manage Models
- ✓ Track Experiments
- ✓ Deploy Models

Custom AI

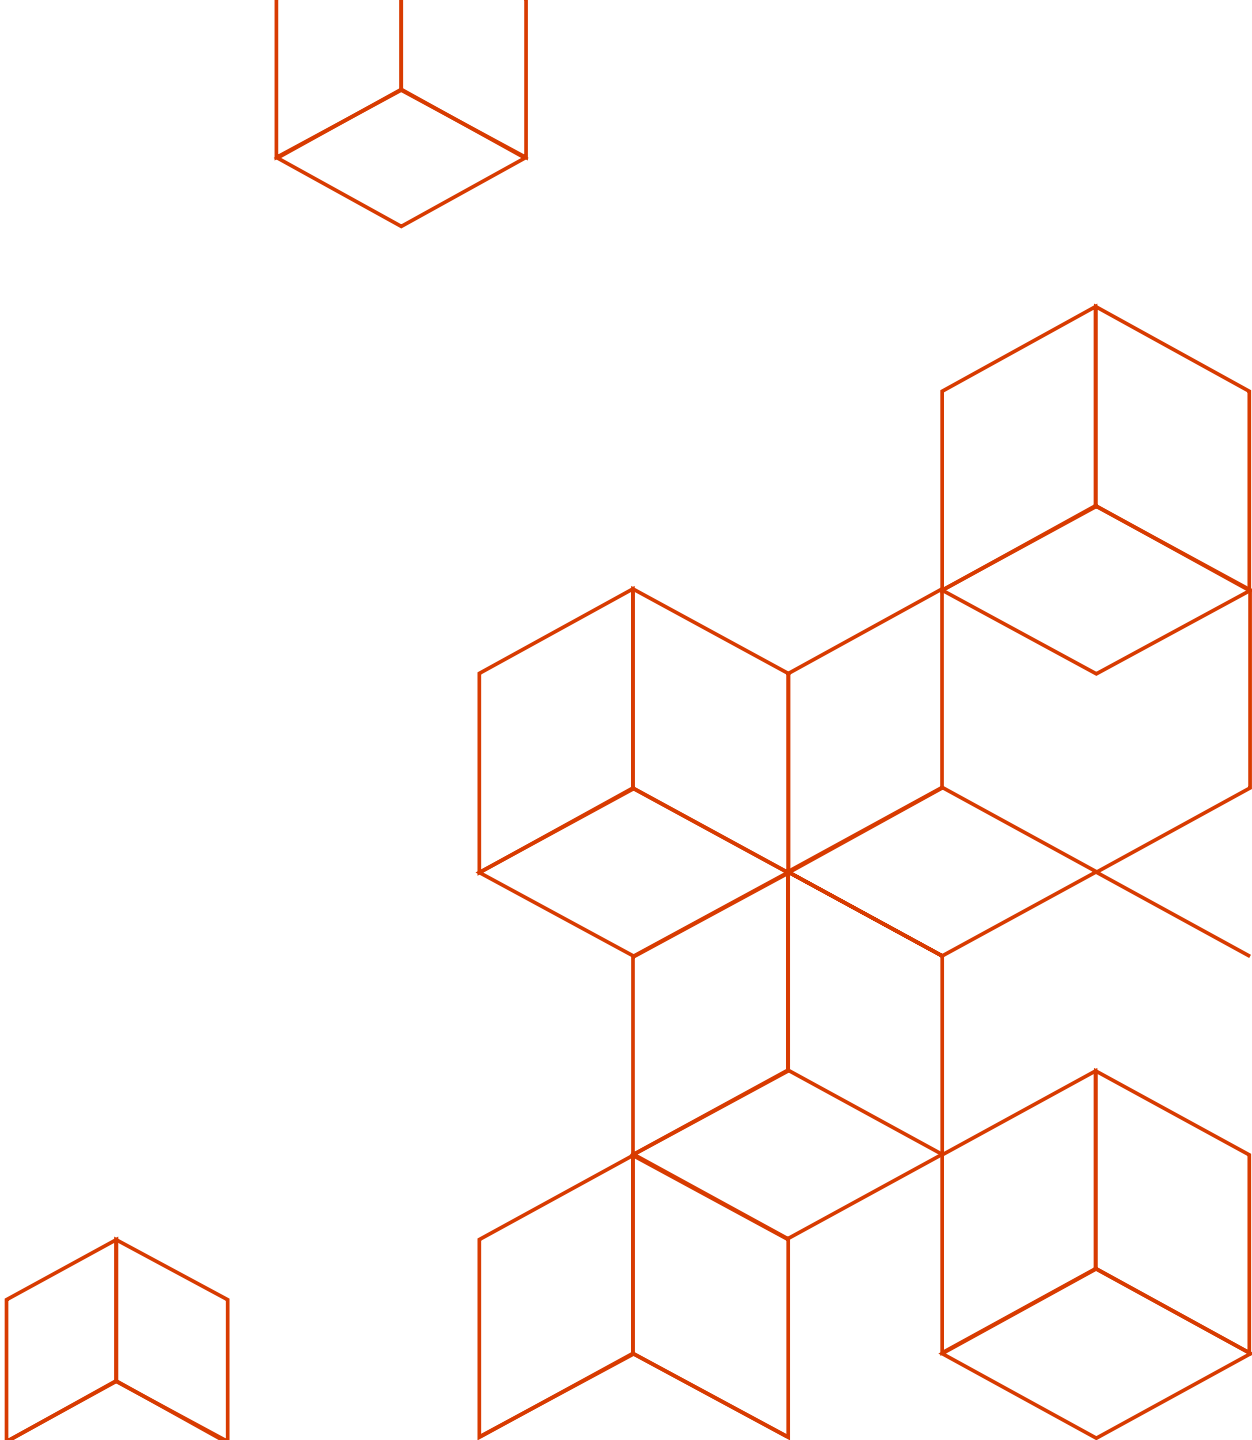

Productive Services

Empower data science and development teams

Integrated data science & data engineering teams

Desktop solutions not adequate

Need a unified big data & machine learning solution

Individual data scientists

Desktop solutions adequate

Need cloud for sporadic compute needs

Azure Databricks + Azure Machine Learning service

Machine Learning VMs

Azure Machine Learning Service

Bring AI to everyone with an end-to-end, scalable, trusted platform

Boost your data science productivity

Increase your rate of experimentation

Deploy and manage your models everywhere

Built with your needs in mind

- Automated machine learning
- Managed compute
- DevOps for machine learning
- Simple deployment
- Tool agnostic Python SDK
- Support for open source frameworks

Seamlessly integrated with the Azure Portfolio

Data Science Virtual Machine

Data Science Virtual Machines (DSVM)

Pre-Configured environments in the cloud for
Data Science & AI Modeling,
Development & Deployment.

Samples to get started

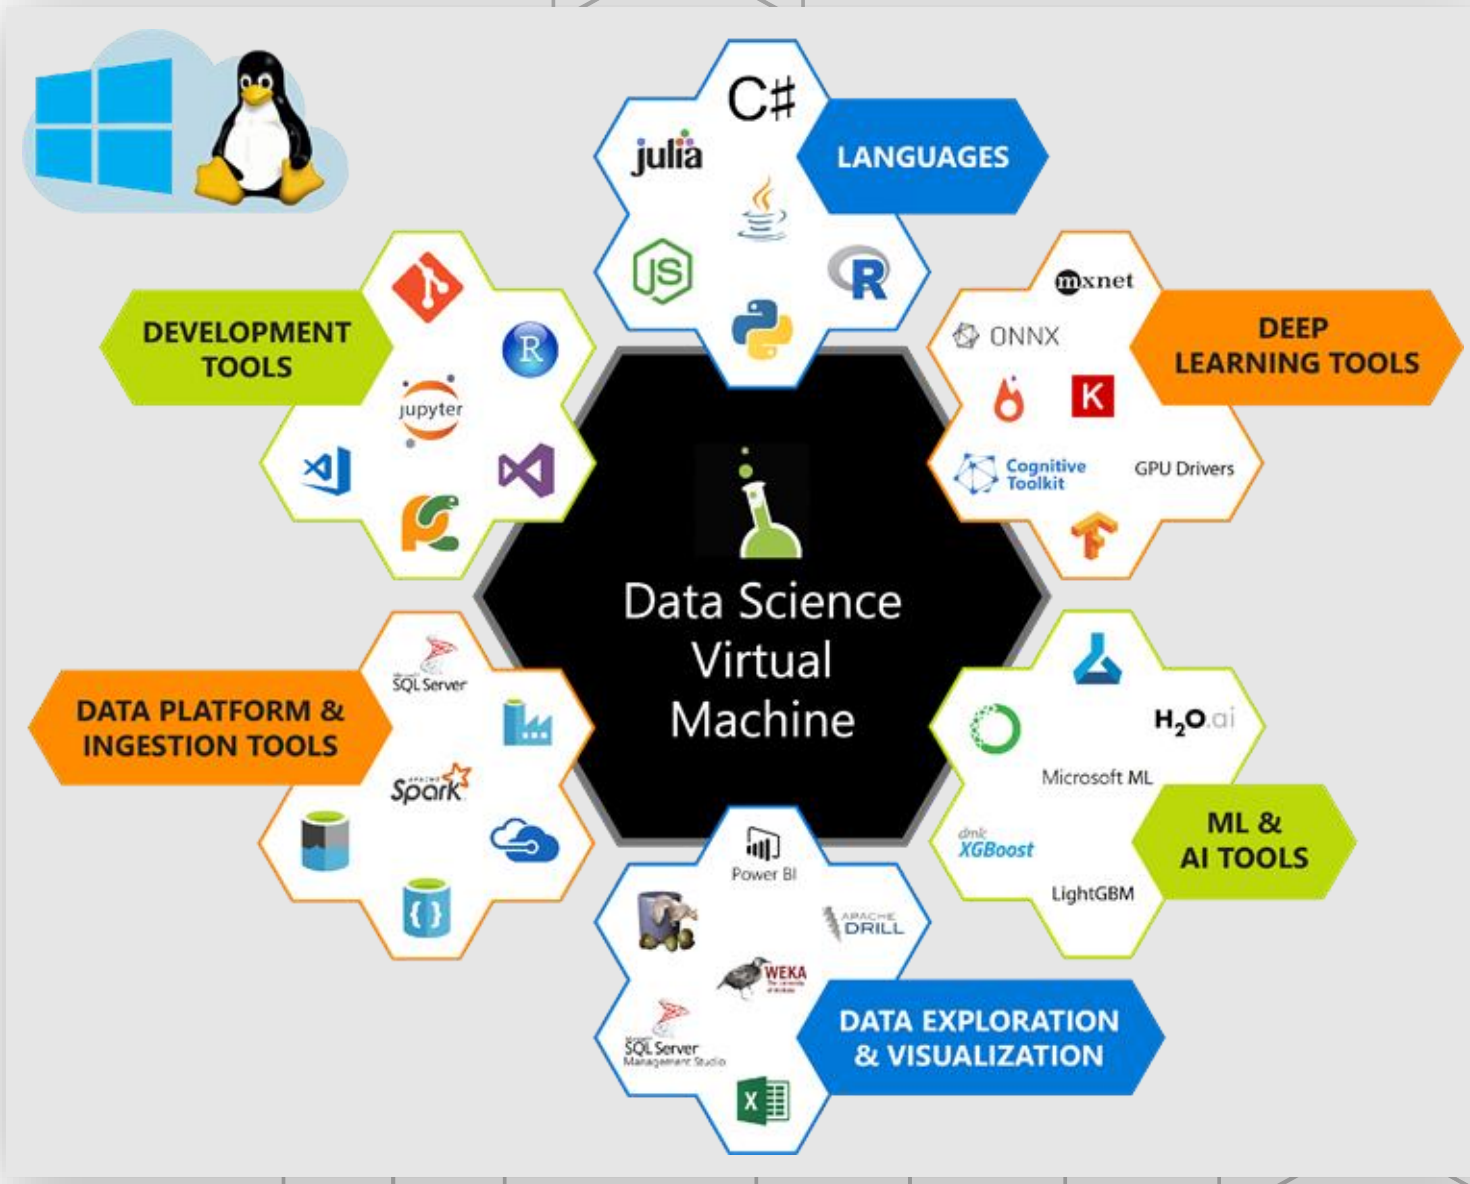

Why Data Science VMs?

Elastic analytics desktop in the cloud

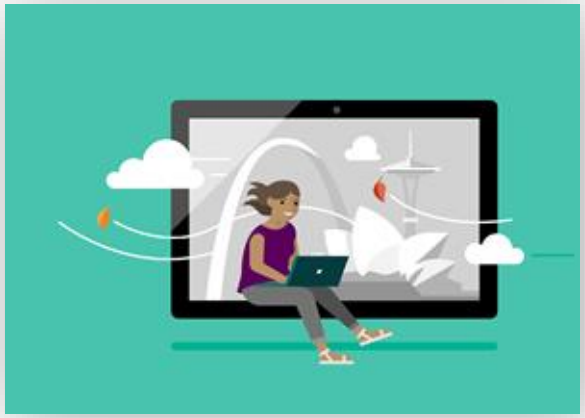

Get started quickly on Azure Machine Learning

Pre-setup environment at work and school

Enterprise Ready and globally available

Managed Compute

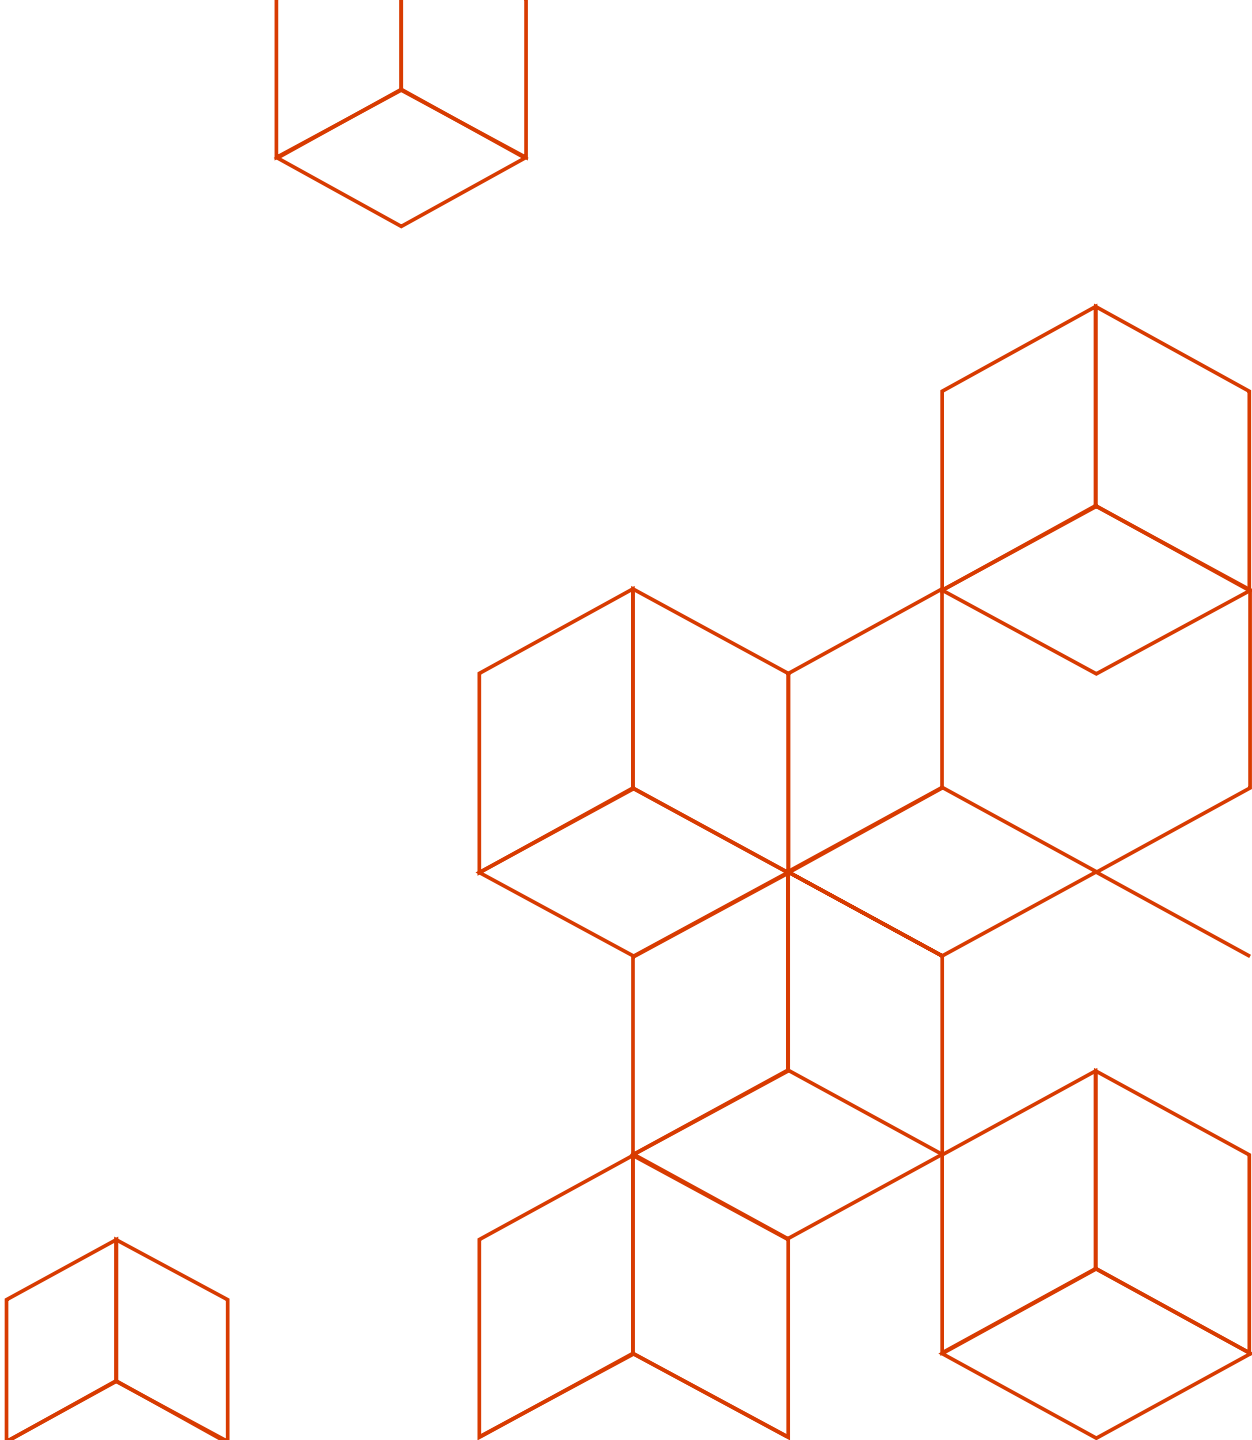

Distributed training on managed compute

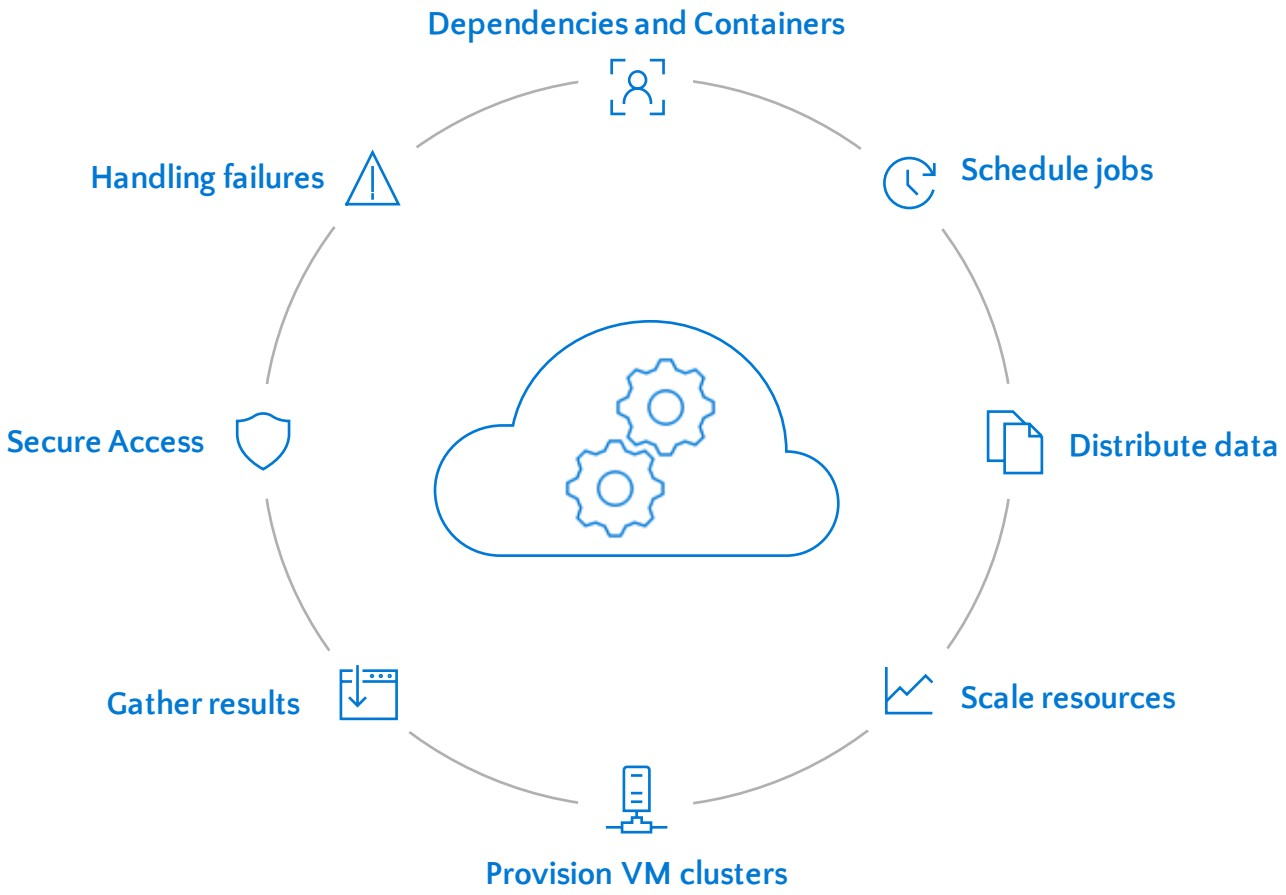

Training Infrastructure

Dependencies and Containers

Leverage system-managed AML compute or bring your own compute

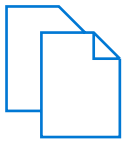

Distribute data

Manage and share resources across a workspace

Schedule jobs

Train at cloud scale using a framework of choice

Scale resources

Autoscale resources to only pay while running a job

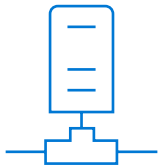

Provision VM clusters

Use the latest NDv2 series VMs with the NVIDIA V100 GPUs

Powerful Infrastructure

Accelerate deep learning

CPUs

General purpose machine learning

D, F, L, M, H Series

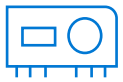

GPUs

Deep learning

N Series

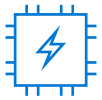

FPGAs

Specialized hardware accelerated deep learning

Project Brainwave

FPGA NEW UPDATES:

Support for image classification and recognition scenarios

ResNet 50, ResNet 152, VGG-16, SSD-VGG, DenseNet-121

DevOps for Machine Learning

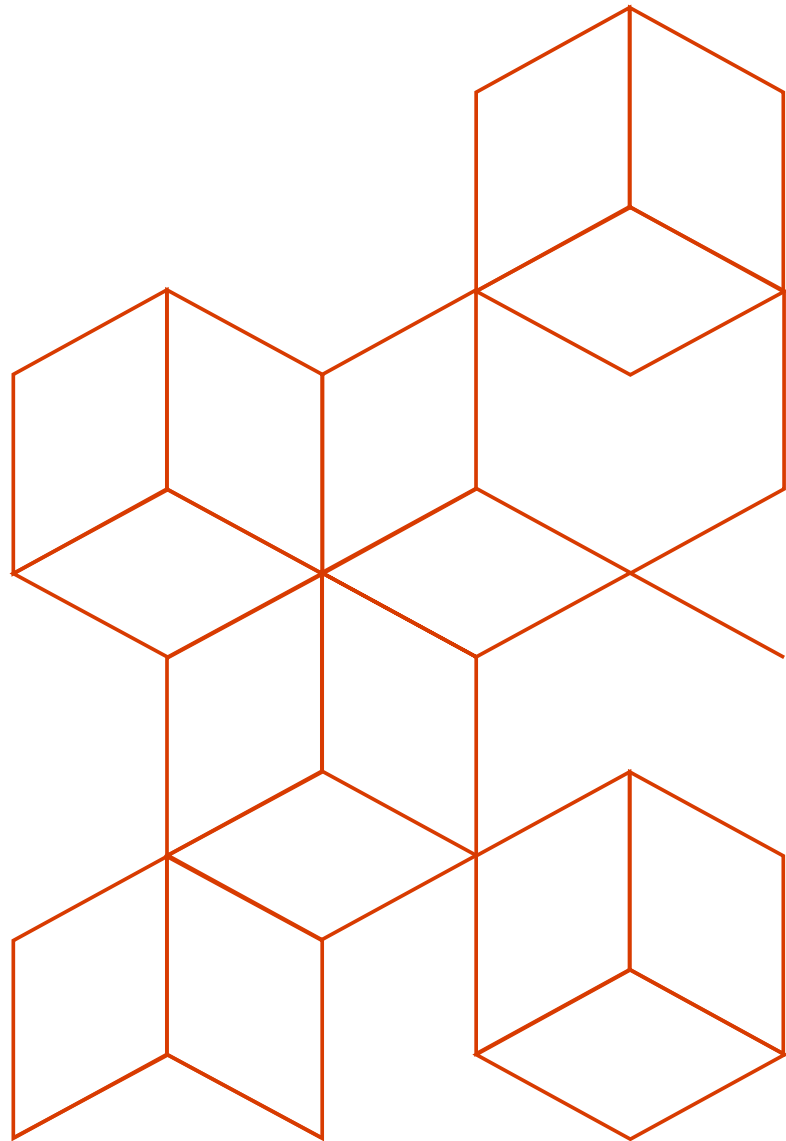

DevOps loop for data science

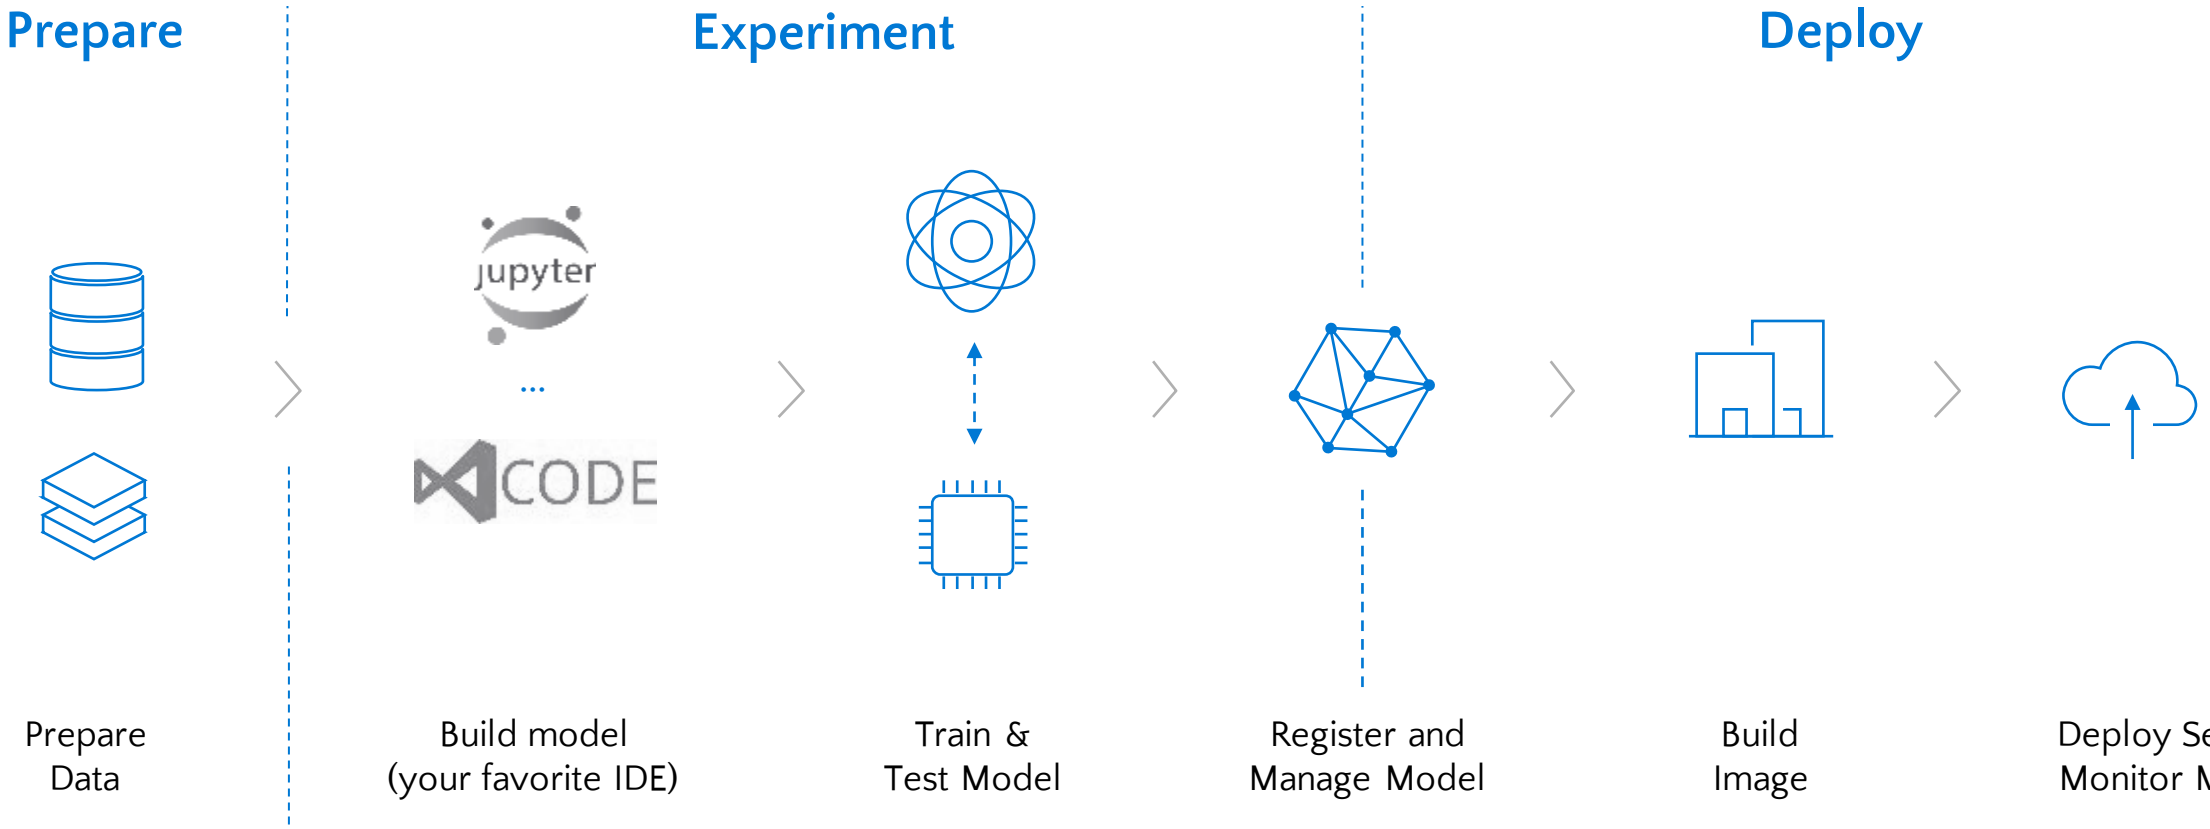

DevOps loop for data science

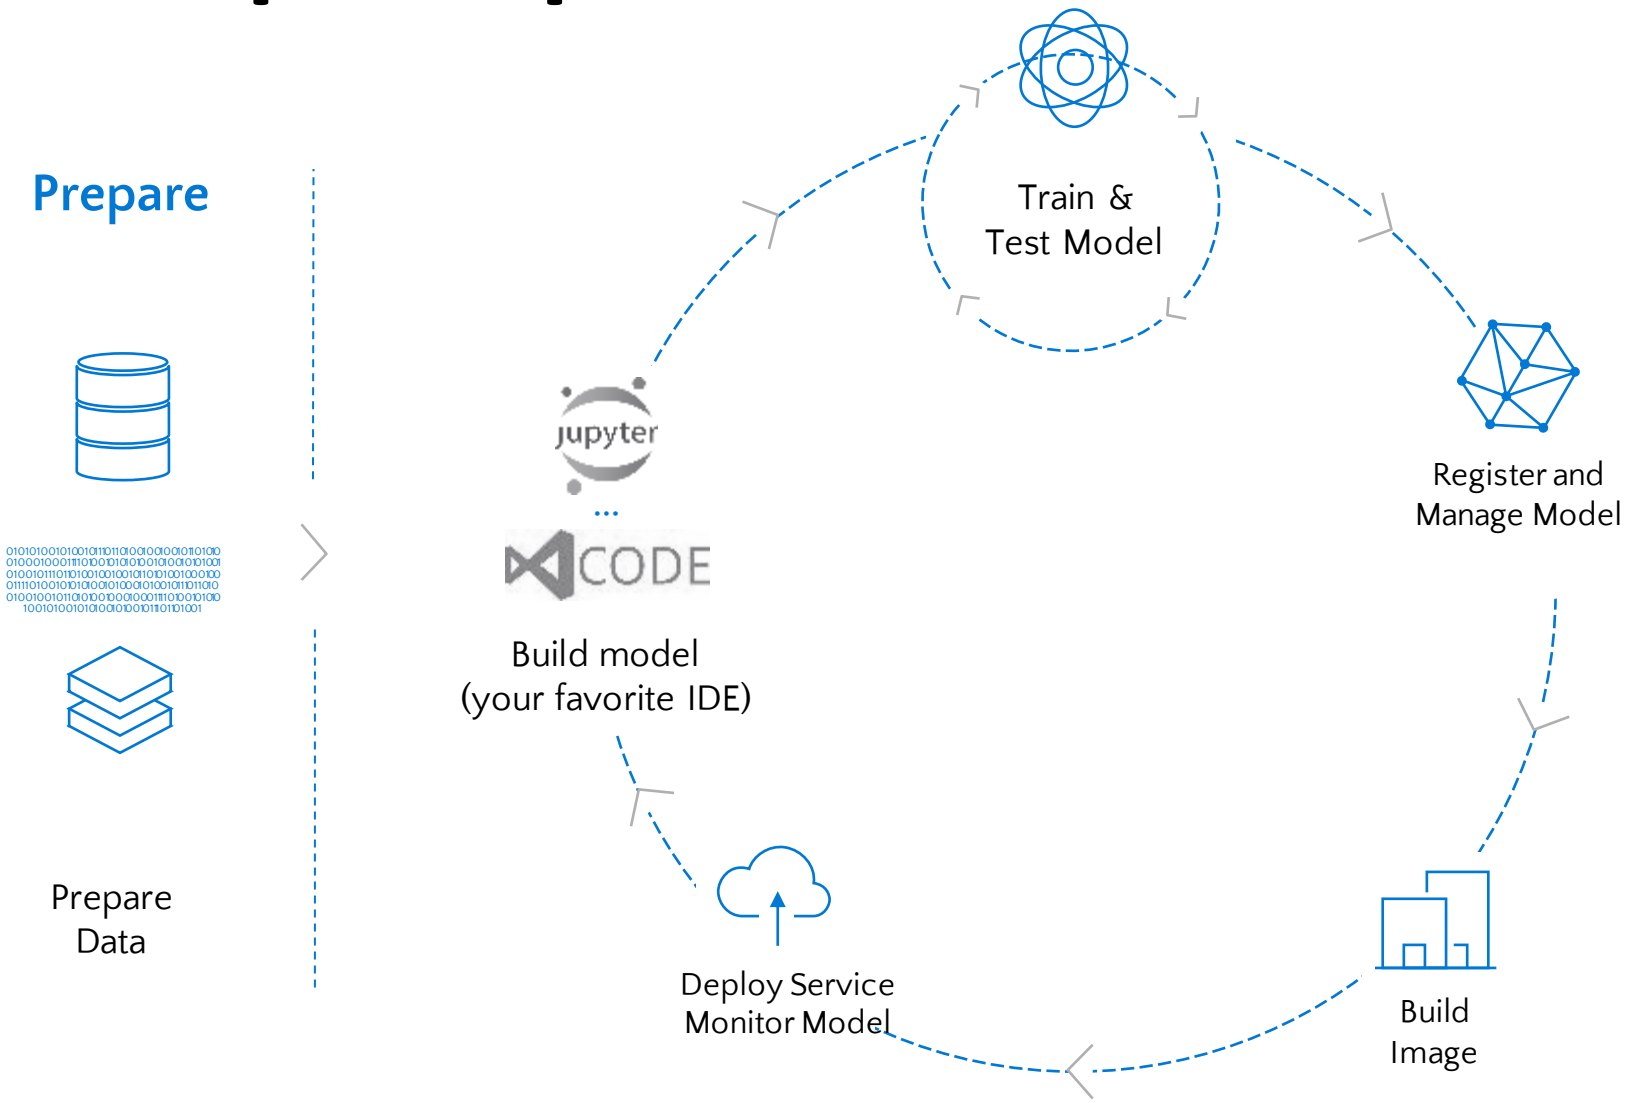

Model Management in Azure Machine Learning

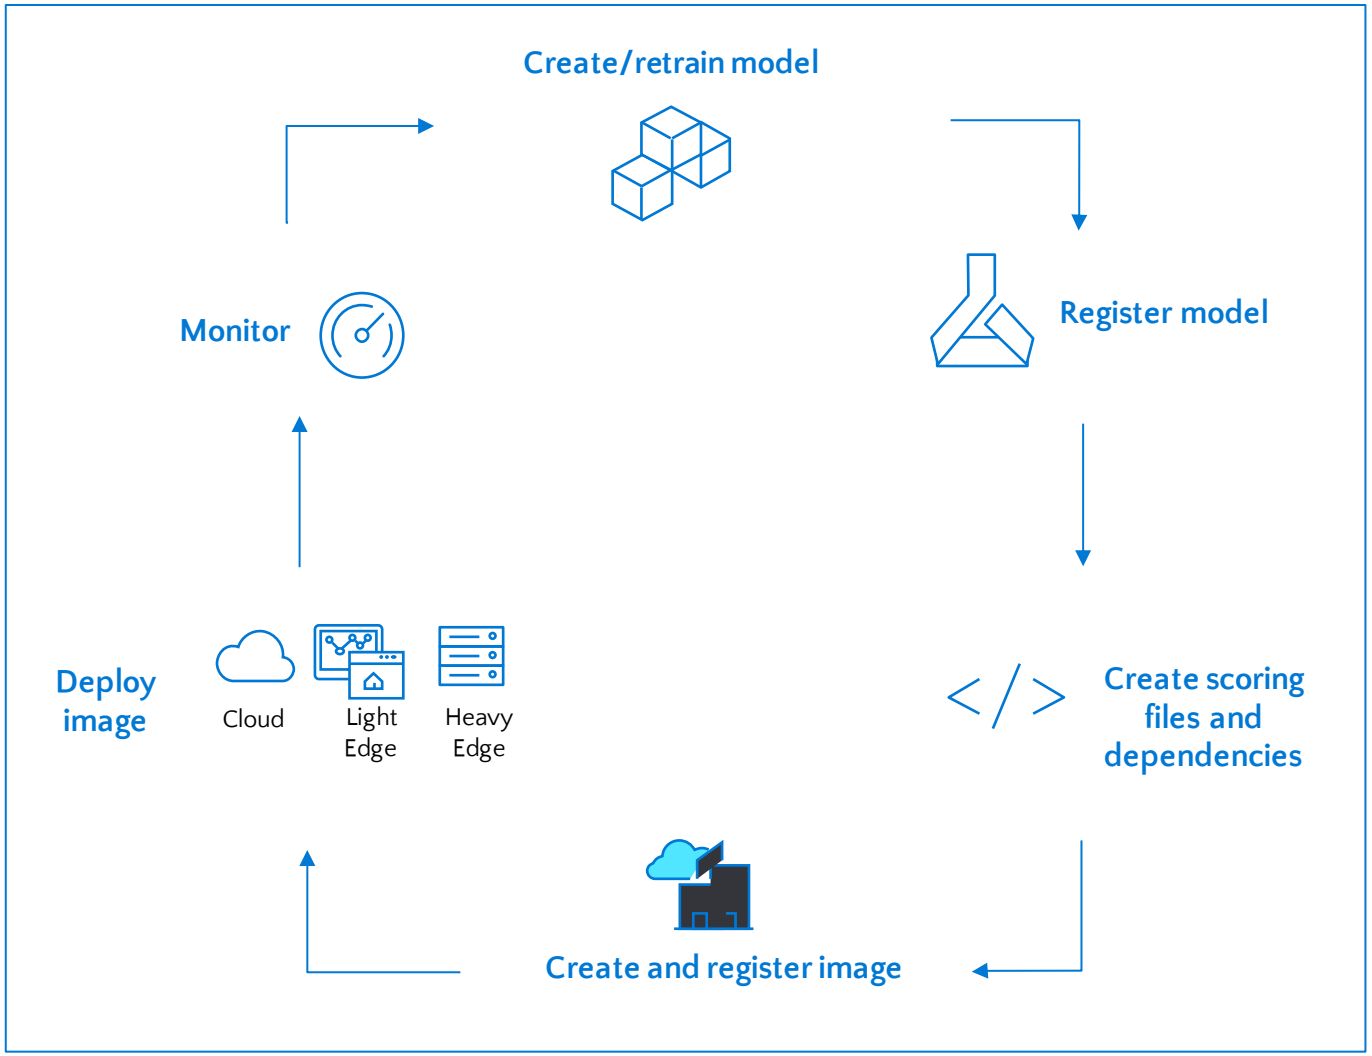

Model Management in detail

Create/Retrain Model

Enable DevOps with full CI/CD integration with VSTS

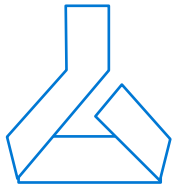

Register Model

Track model versions with a central model registry

Monitor

Oversea deployments through Azure AppInsights

Experimentation

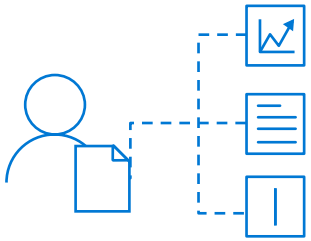

Leverage service-side capture of run metrics, output logs and models

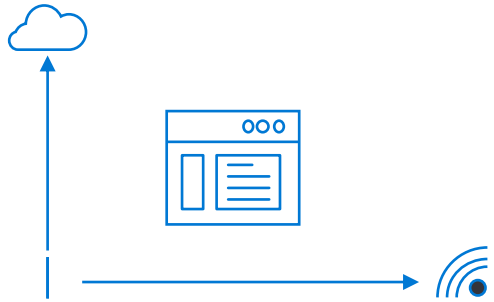

Manage training jobs locally, scaled-up or scaled-out

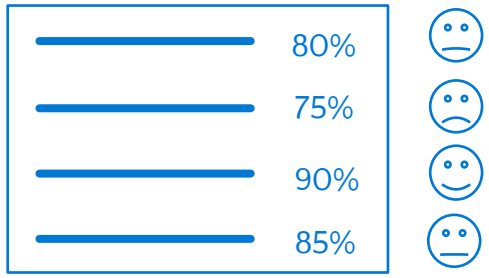

Use leaderboards, side by side run comparison and model selection

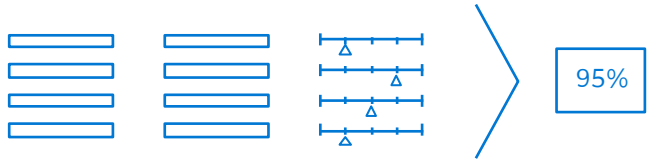

Conduct a hyperparameter search on traditional ML or DNN

Azure Machine Learning Pipelines

Azure Machine Learning Pipelines

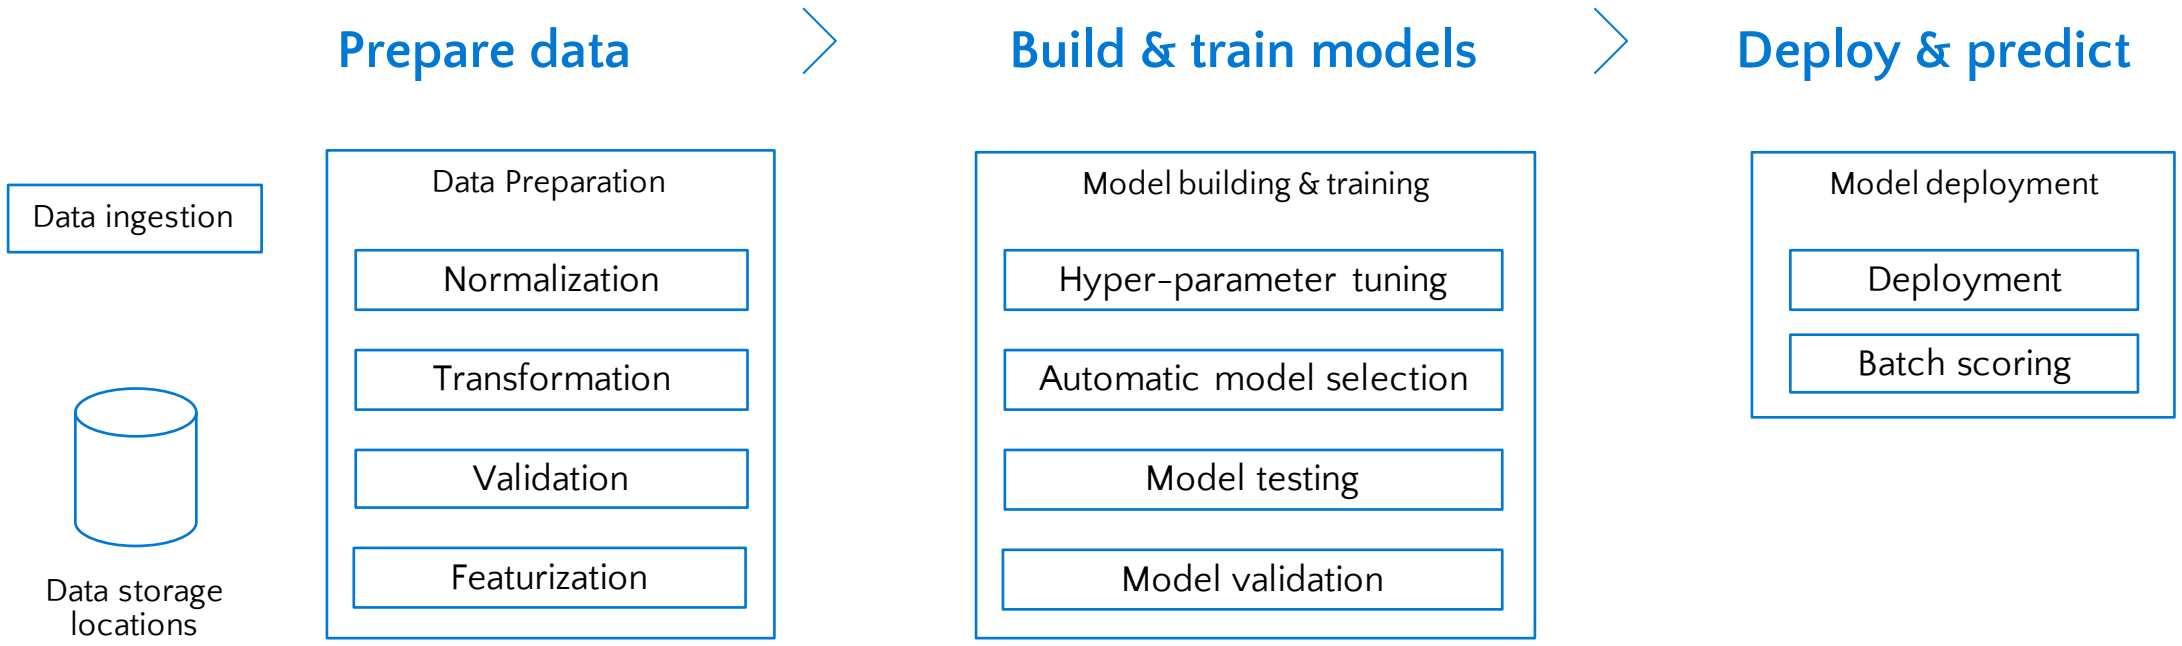

Azure Machine Learning Pipelines

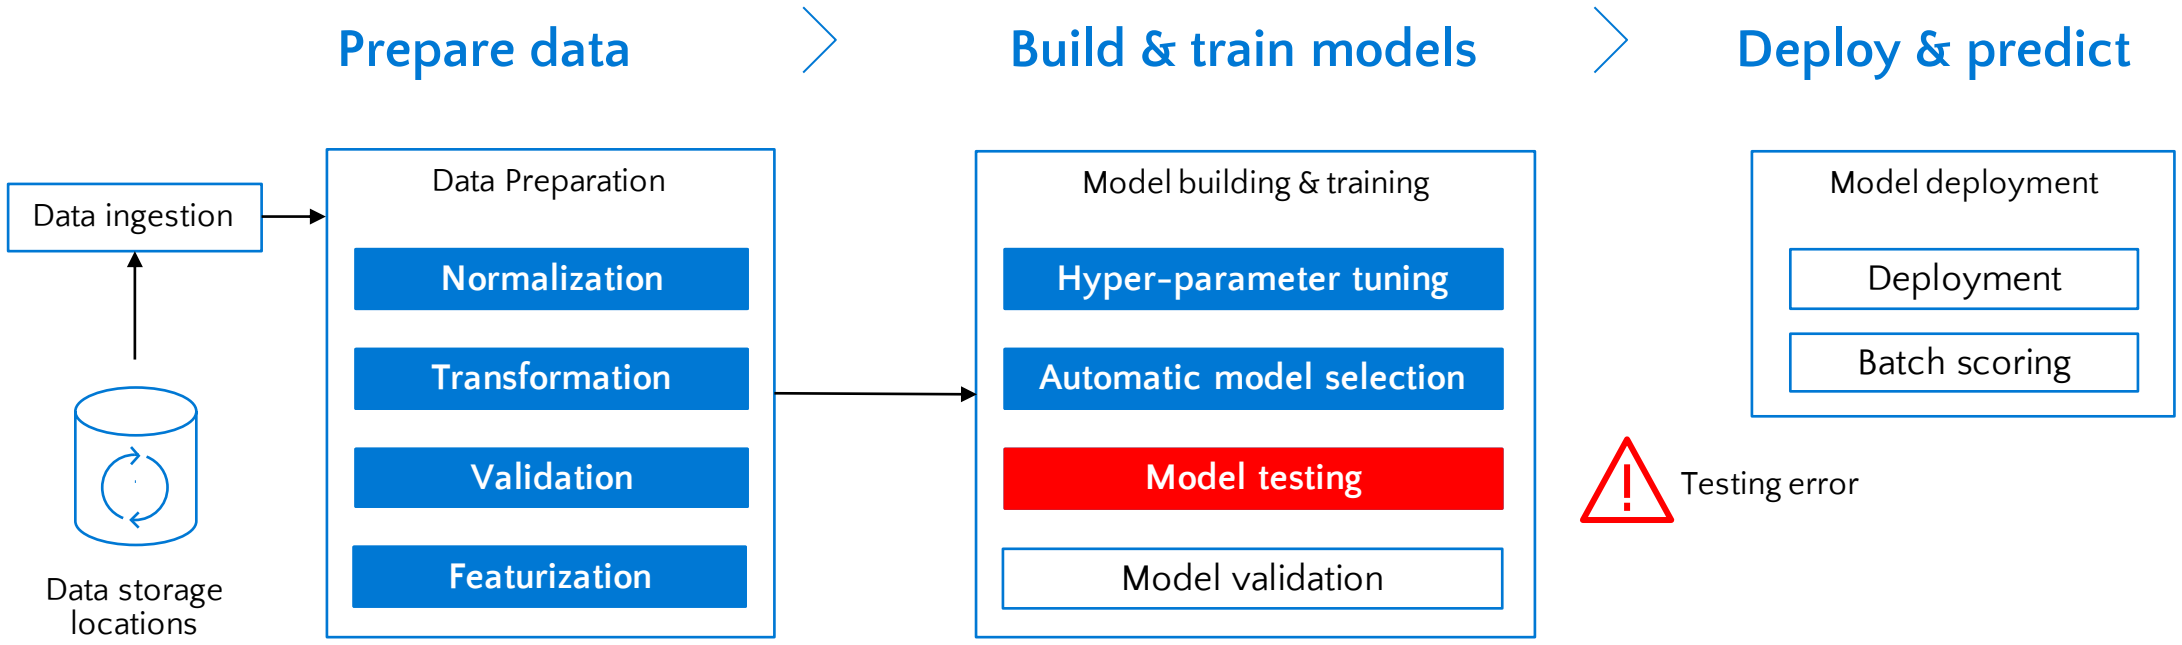

Azure Machine Learning Pipelines

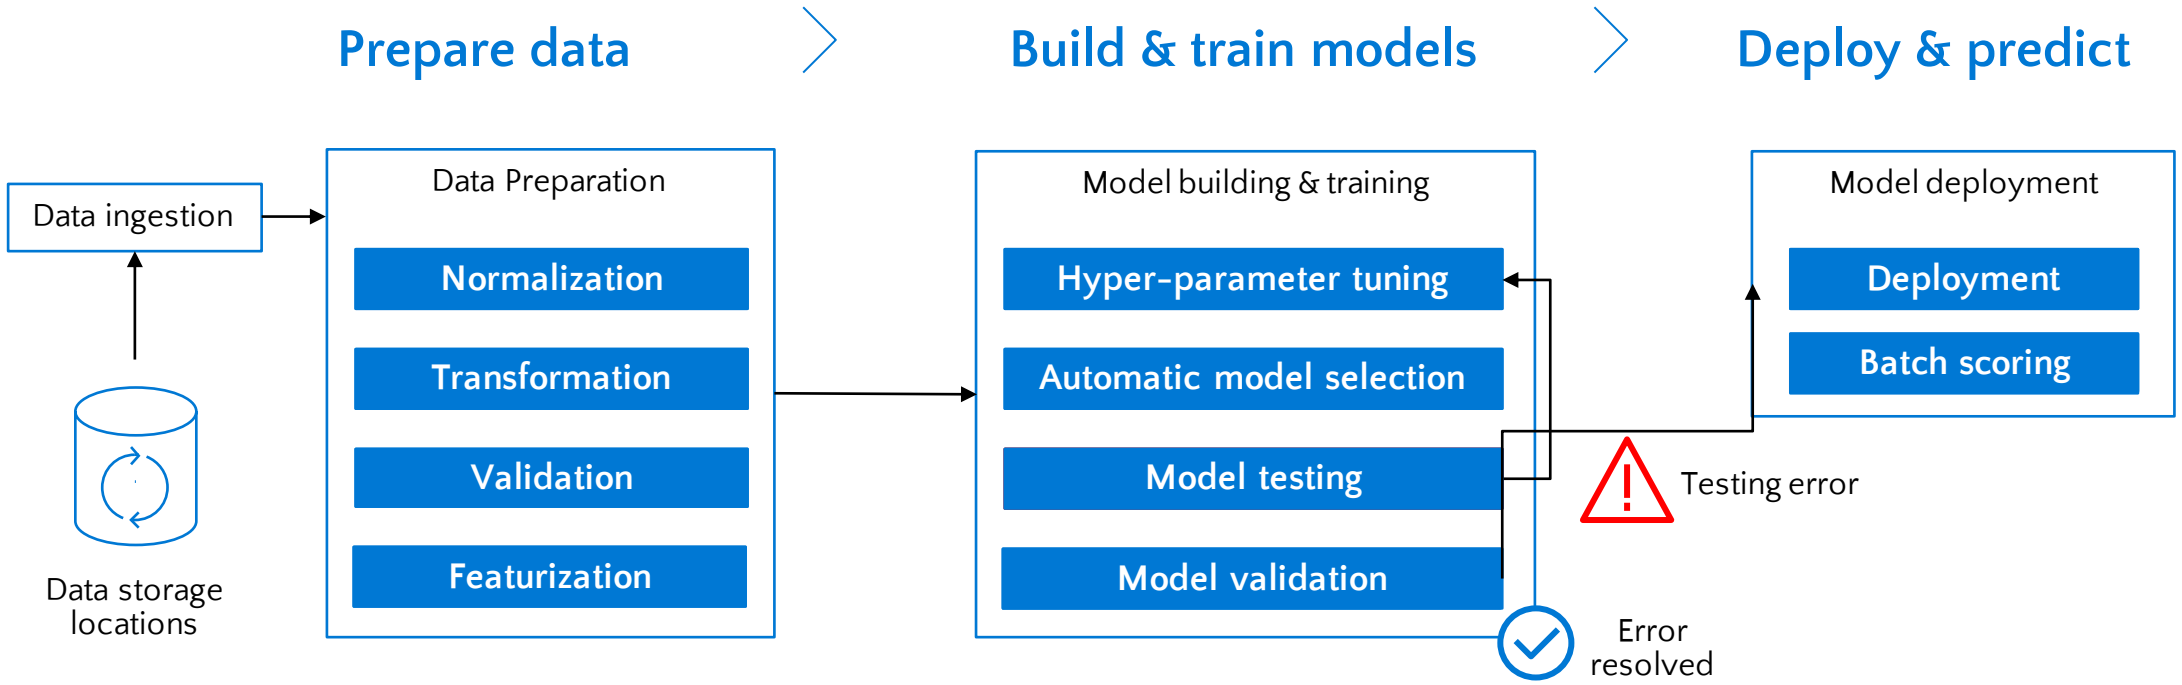

Azure Machine Learning Pipelines with new data

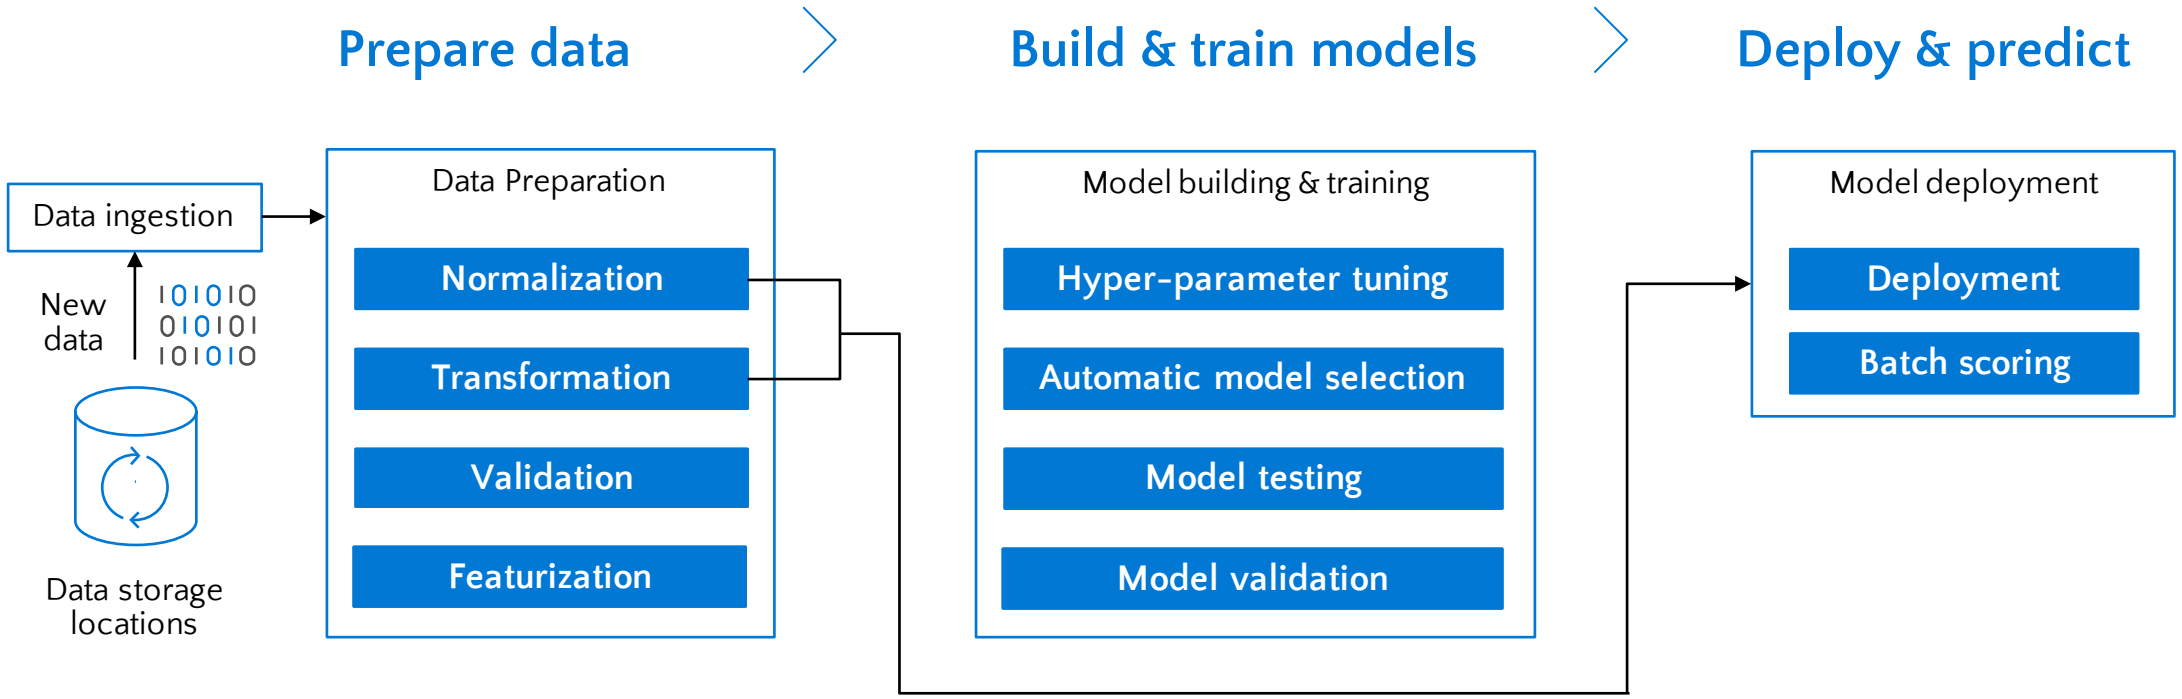

Advantages of Azure ML Pipelines

Unattended runs

Schedule a few steps to run in parallel or in sequence to focus on other tasks while your pipeline runs

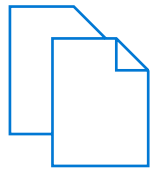

Tracking and versioning

Name and version your data sources, inputs and outputs with the pipelines SDK

Reusability

Create templates of pipelines for specific scenarios such as retraining and batch scoring

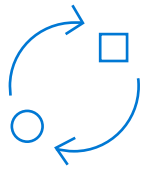

Mixed and diverse compute

Use multiple pipelines that are reliably coordinated across heterogeneous and scalable computes and storages

Support for
Open source frameworks

Popular Frameworks

Use your favorite machine learning frameworks


TensorFlow


PyTorch


Scikit-Learn


MXNet


Chainer


Keras

without getting locked into one framework


ONNX

Community project created by Facebook and Microsoft

Use the best tool for the job. Train in one framework and transfer to another for inference










Create

Frameworks

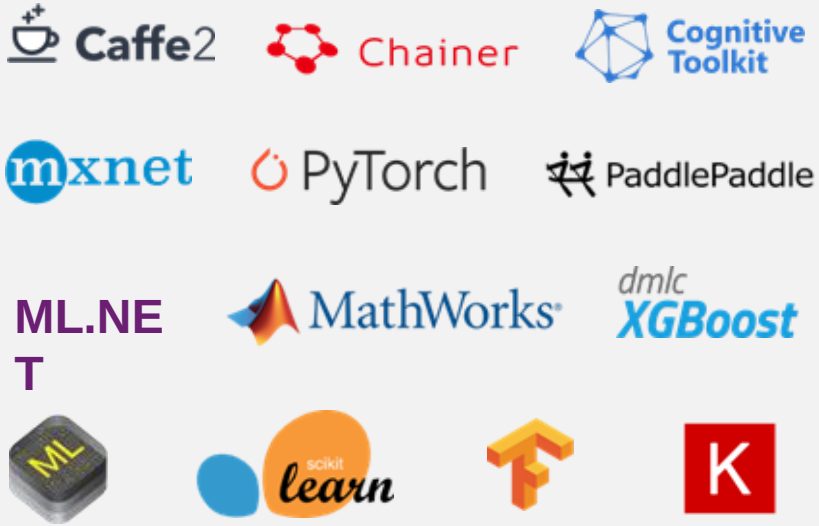

Services

Azure Custom Vision Service

Native support

Converters

Native support

ONNX Model

Native support

Azure

- Azure Machine Learning services
- Ubuntu VM
- Windows Server 2019 VM

Windows Devices

Converters

Other Devices
(iOS, etc)

Deploy

Deployment

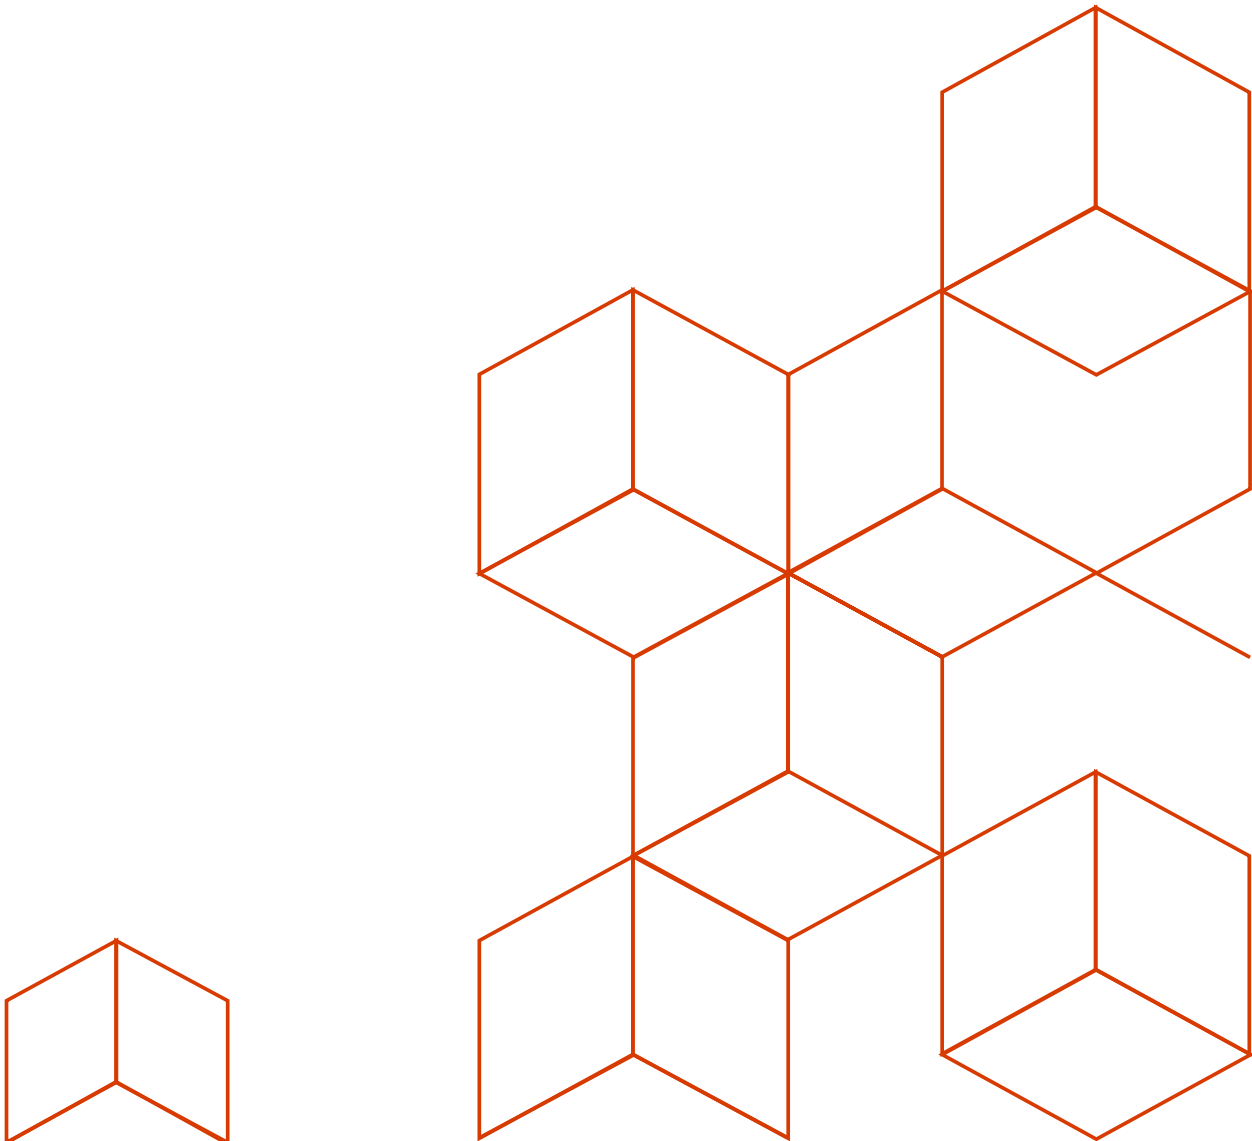

Flexible Deployment

Deploy and manage models on intelligent cloud and edge

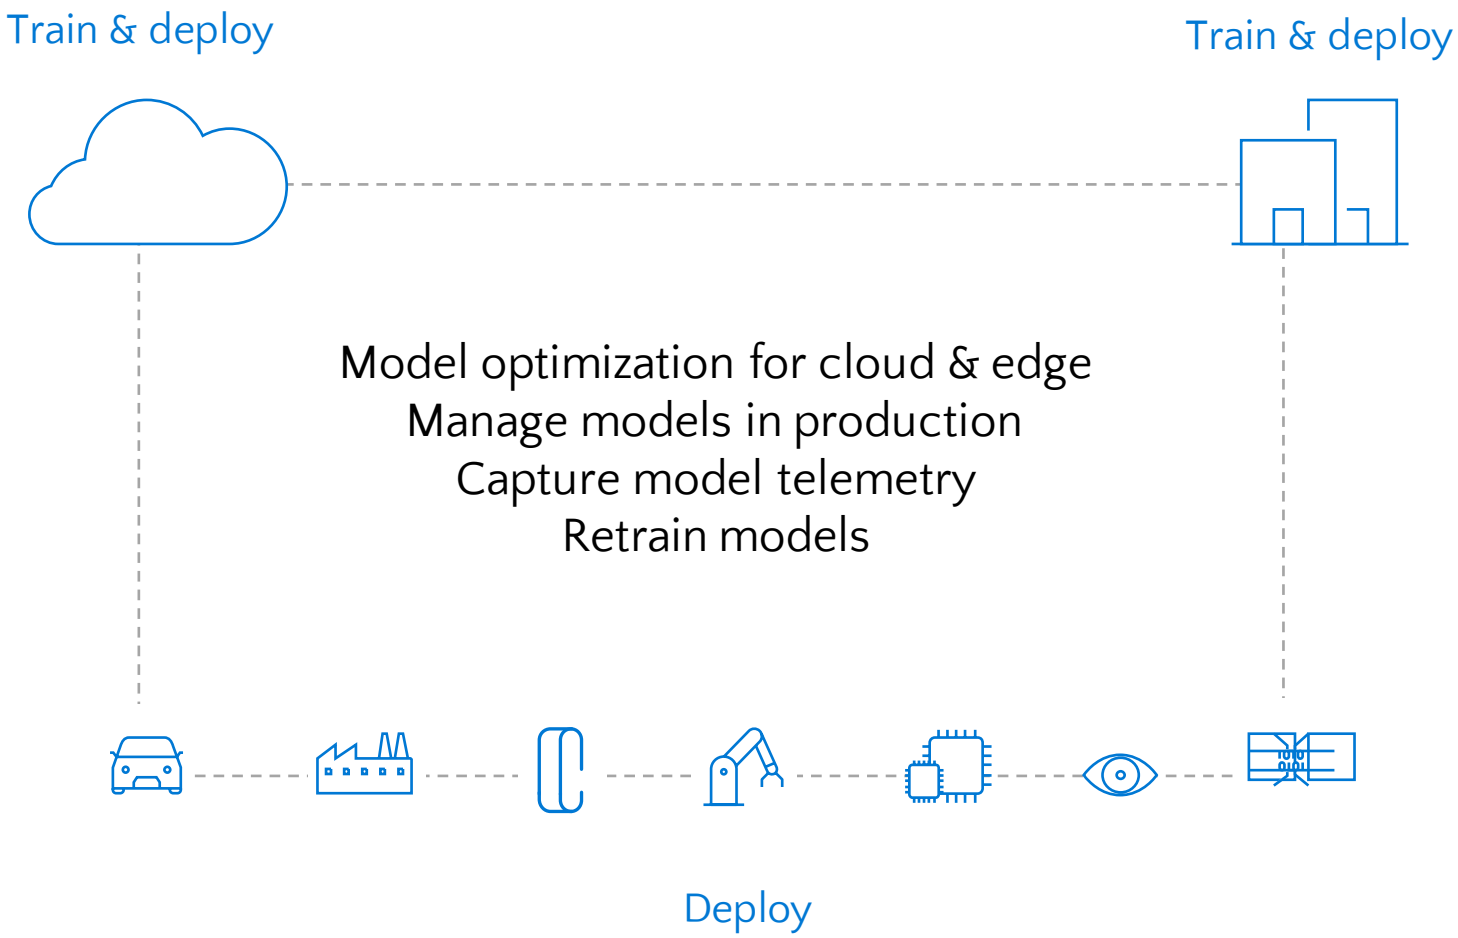

Deploy Azure ML models at scale

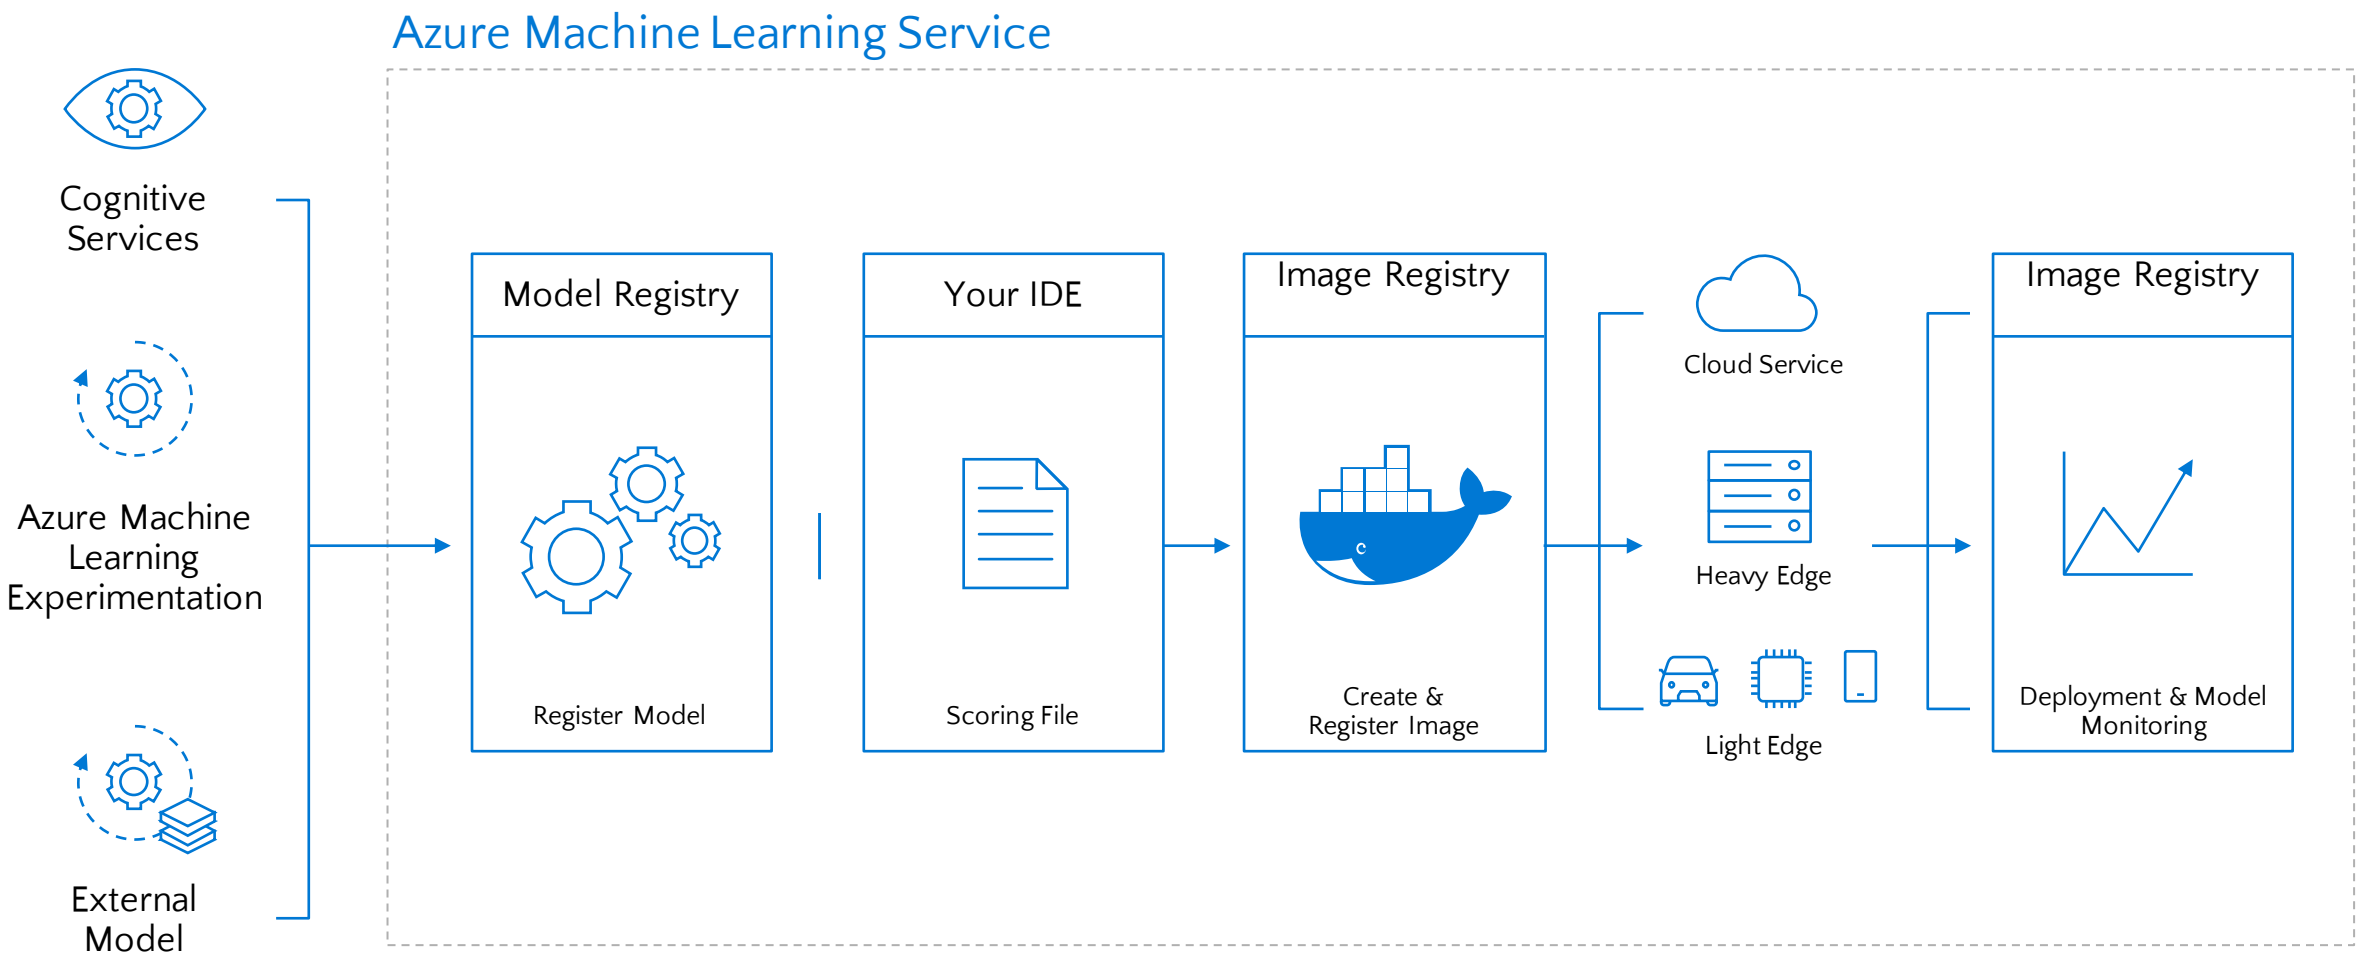

Deployments to Compute Targets

Compute target	Deployment type	Description
Azure Container Instances (ACI)	Web service	Fast deployment. Good for development or testing.
Azure Kubernetes Service (AKS)	Web service	Good for high-scale production deployments. Provides autoscaling, and fast response times.
Azure IoT Edge	IoT module	Deploy models on IoT devices. Inferencing happens on the device.
Field-programmable gate array (FPGA)	Web service	Ultra-low latency for real-time inferencing.

Tool Agnostic Python SDK

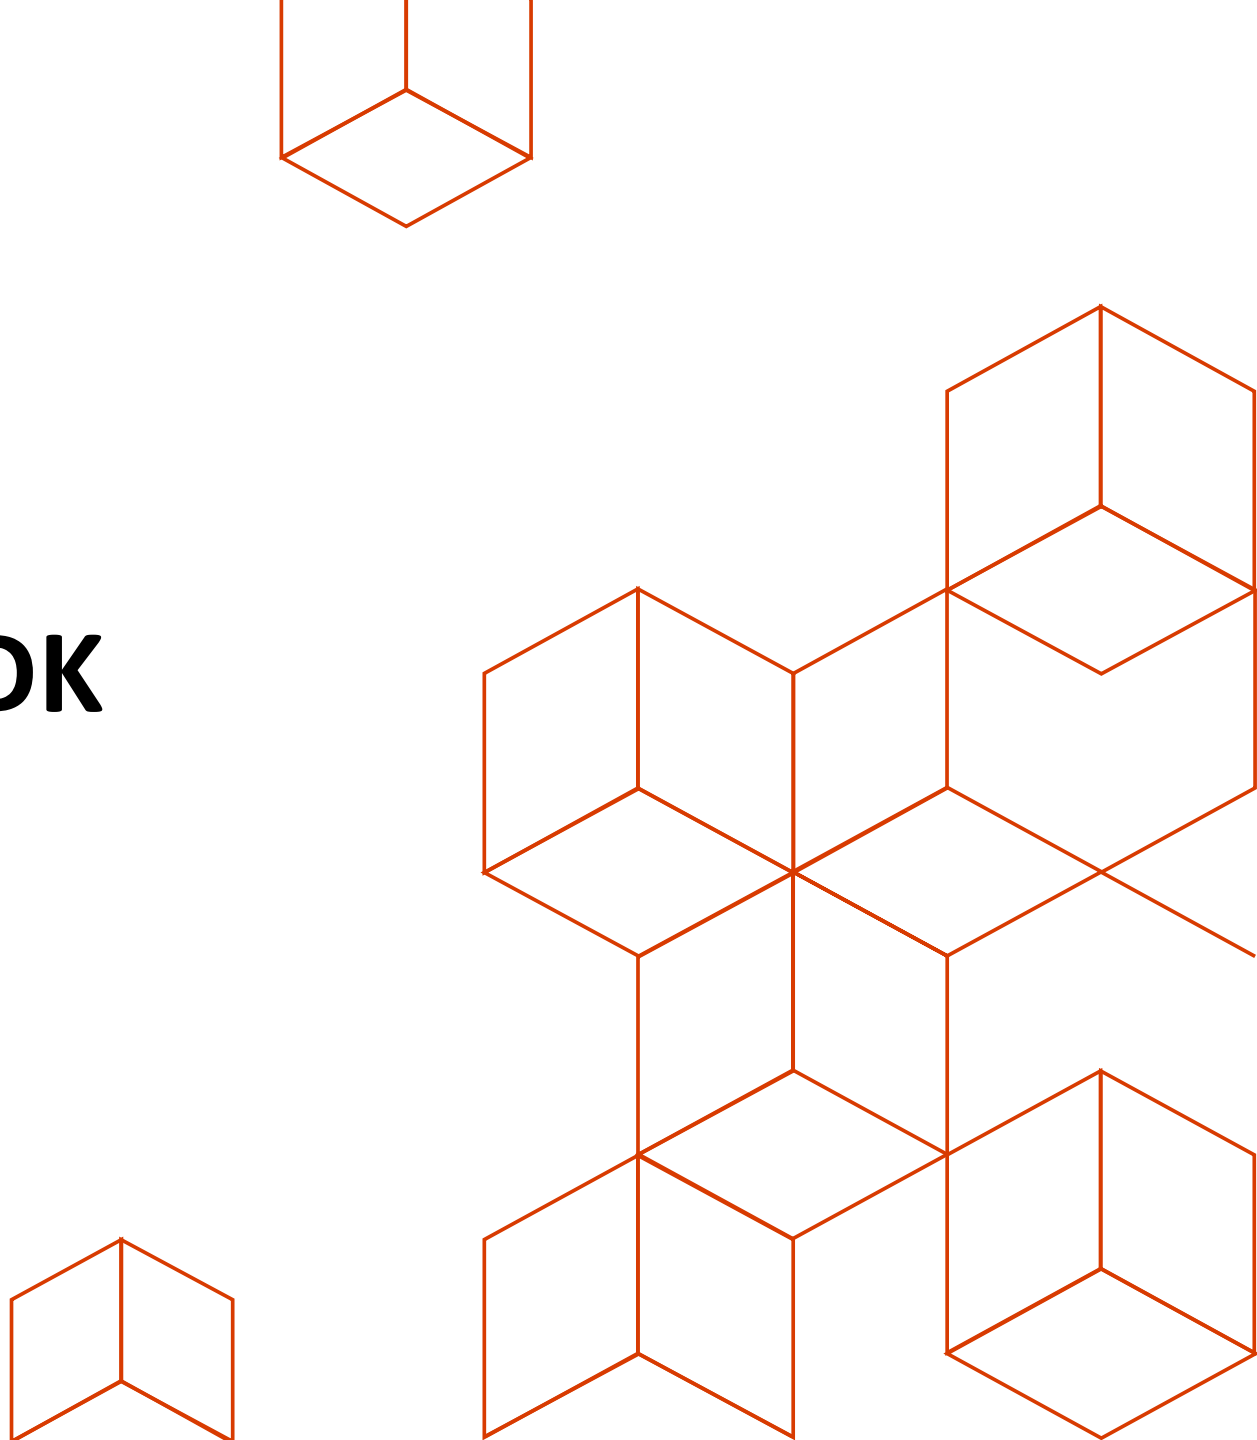

Tool Agnostic Python SDK

Use your favorite IDEs, editors, notebooks, and frameworks

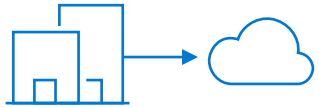

Flexibility of your local environment or curated cloud environment

Integrate with other services like Azure Databricks

Get started quickly without any complex pre-requisites

Azure ML Service includes DataPrep SDK

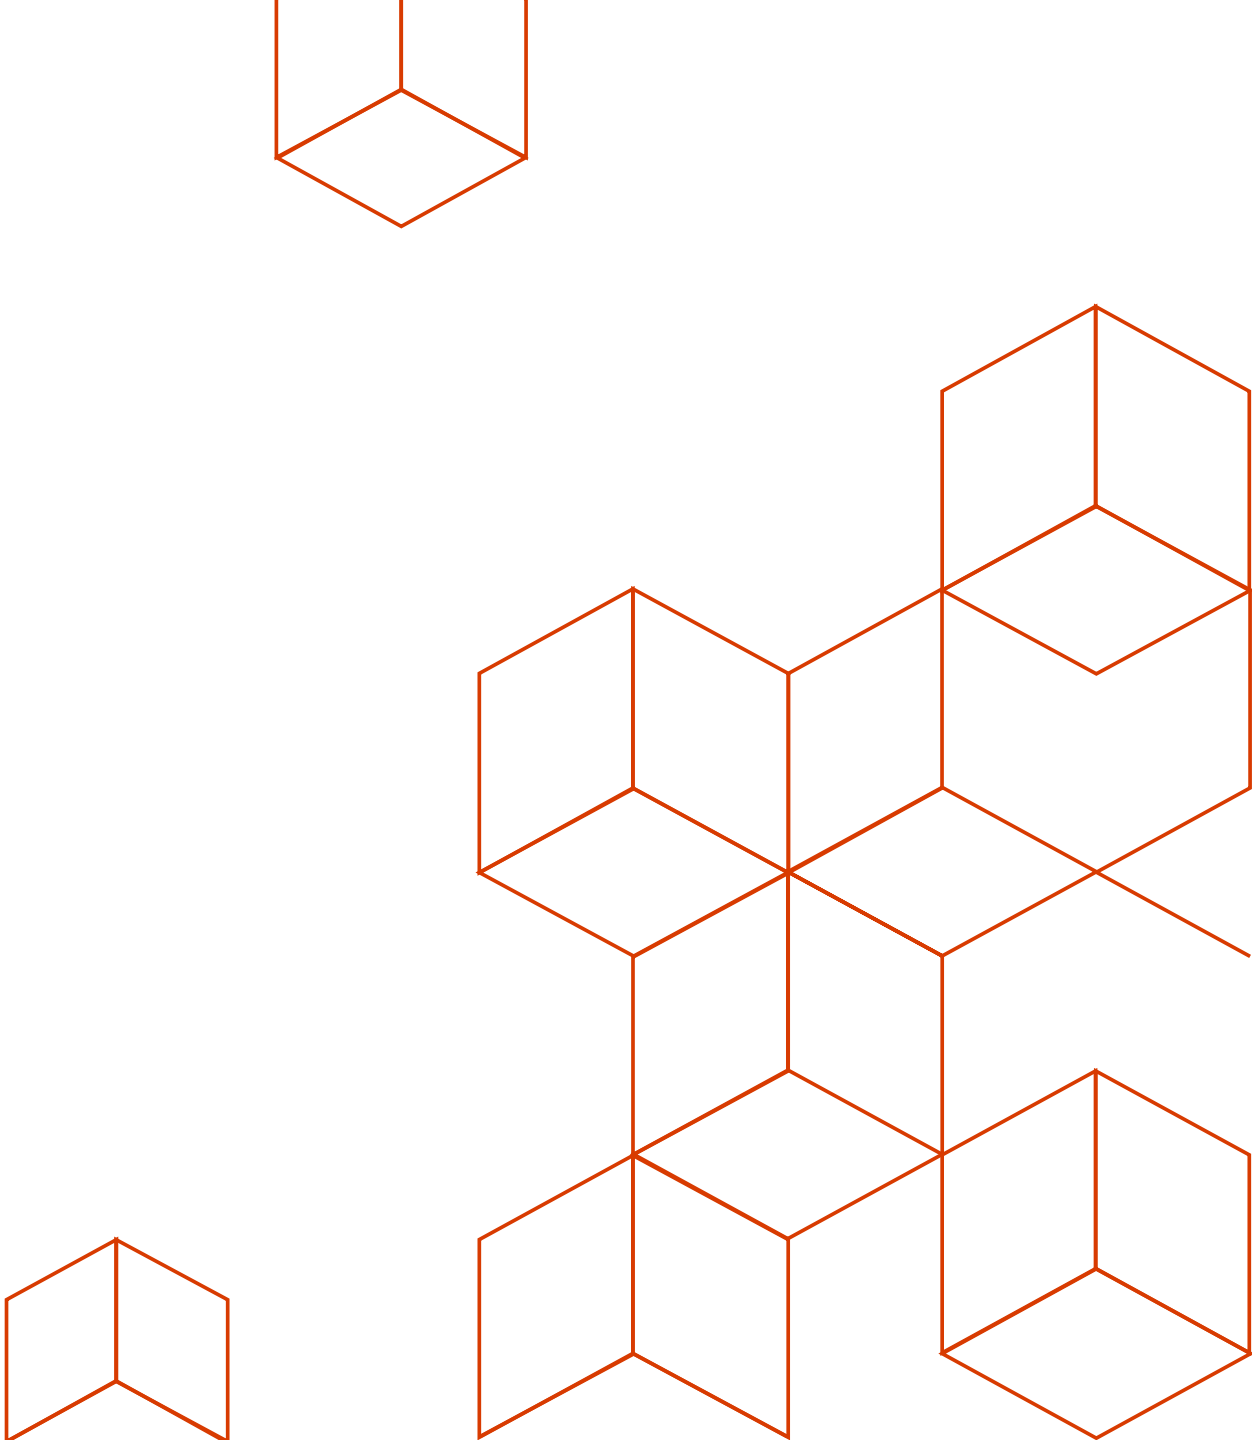

Customer Challenges and Pain Points

- Understanding the semantics of data is difficult and time-consuming
- Merging data from different sources is a manual process
- Detecting, troubleshooting and fixing errors is a high tax
- Custom code is always required
- Operationalization is challenging

Examples of manual, non-scalable work

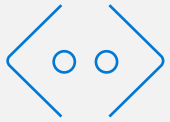

Data formatting

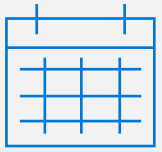

Dealing with dates

'Rectangularizing' data

Data Lifecycle

Interactive training

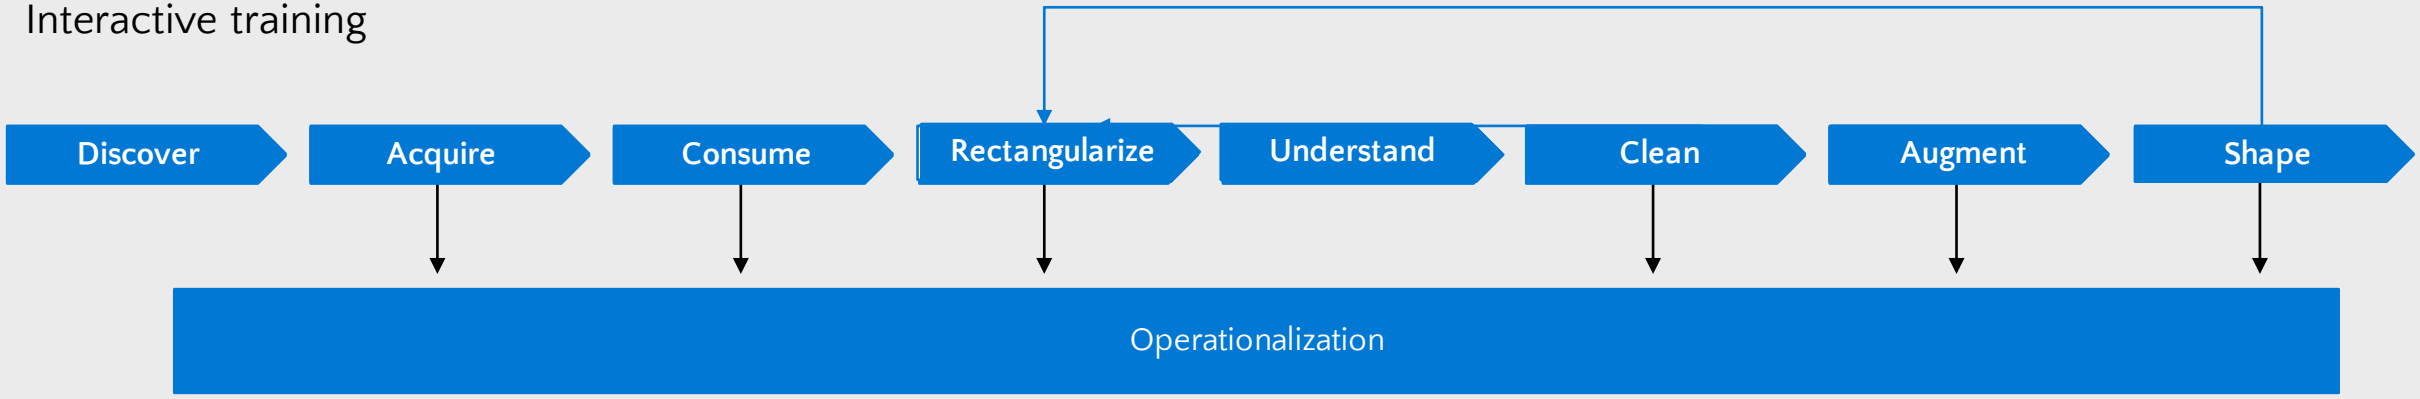

Retraining/scoring

Data Prep SDK

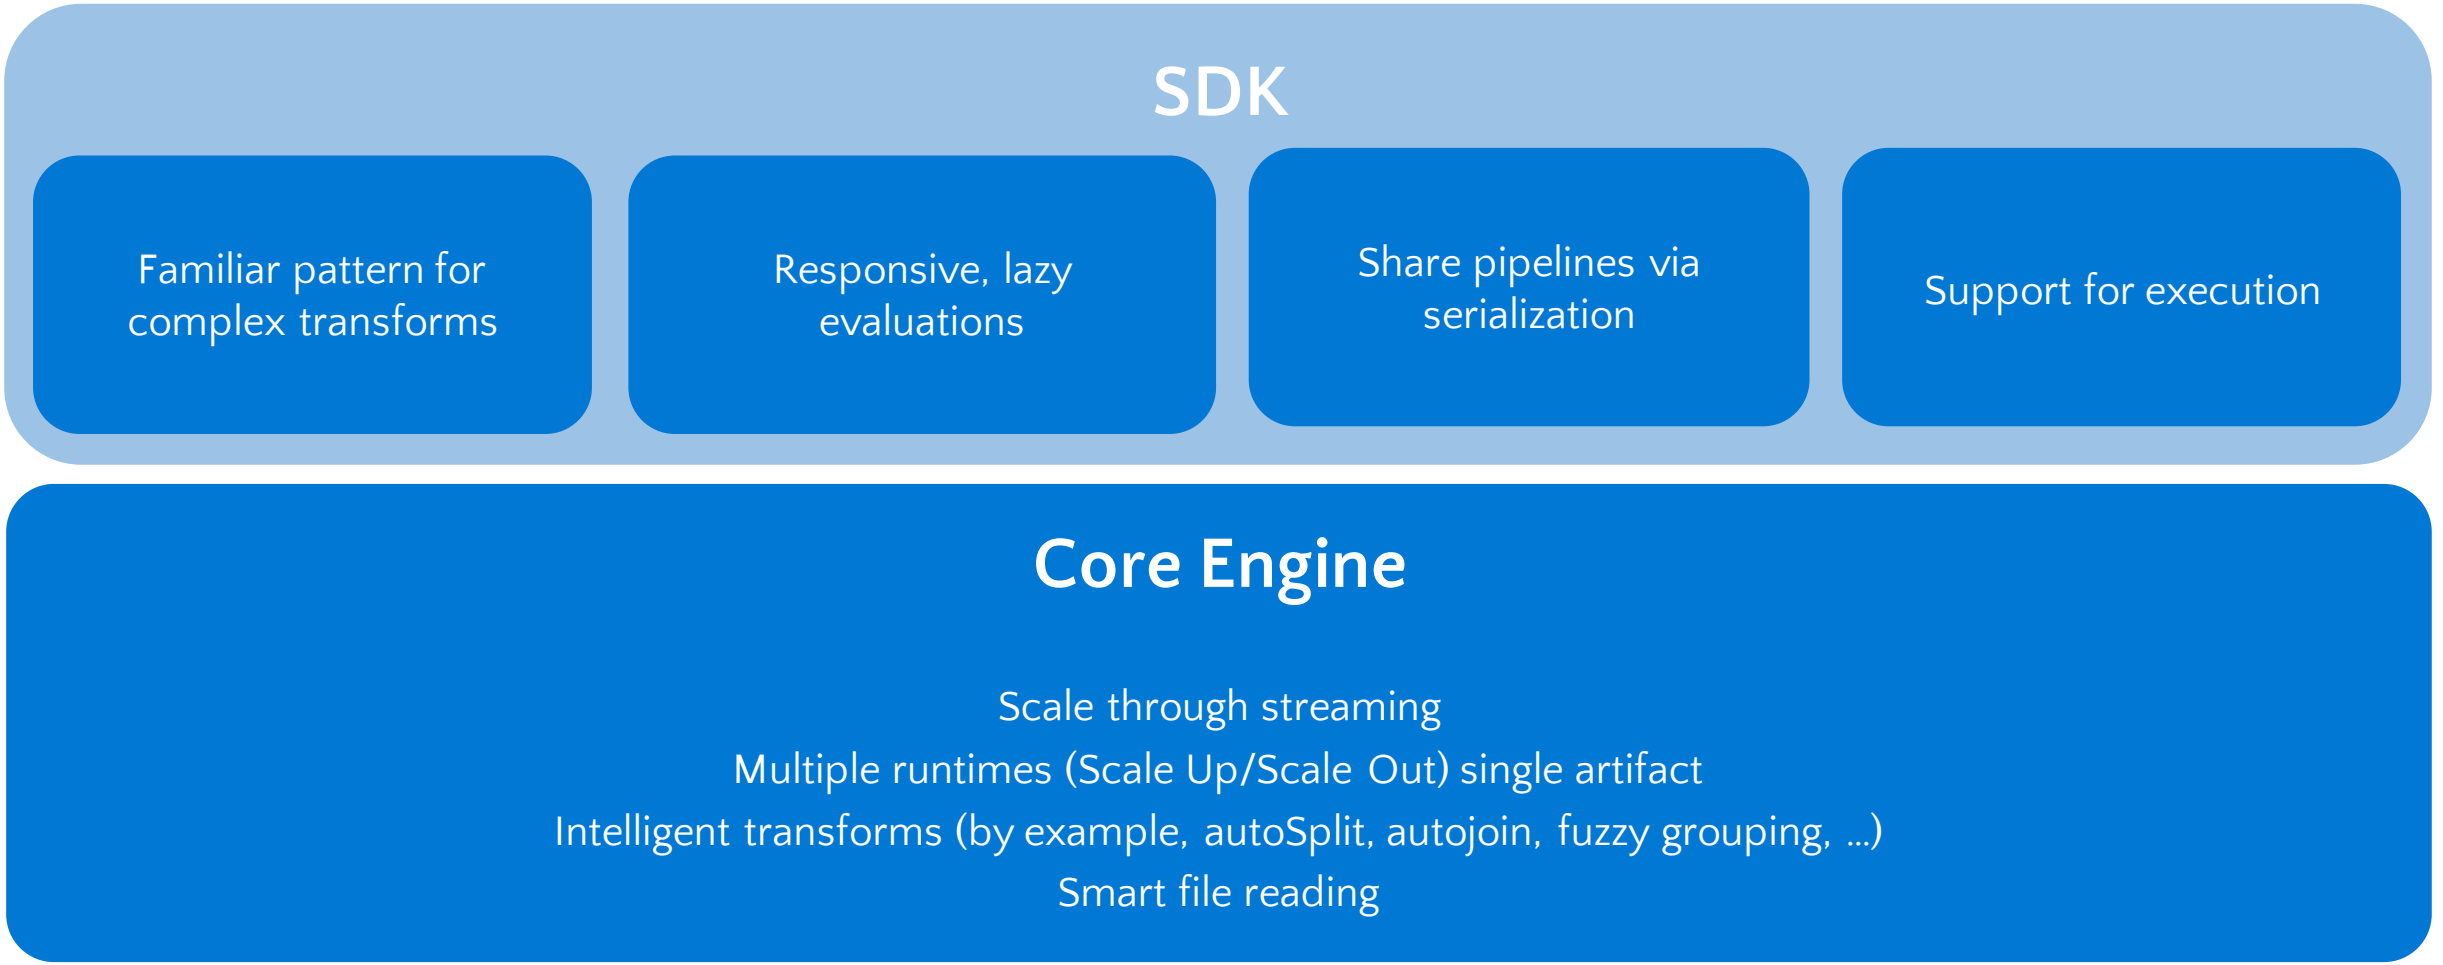

Summary

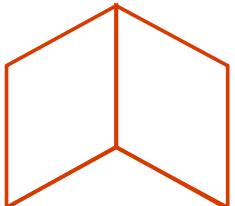

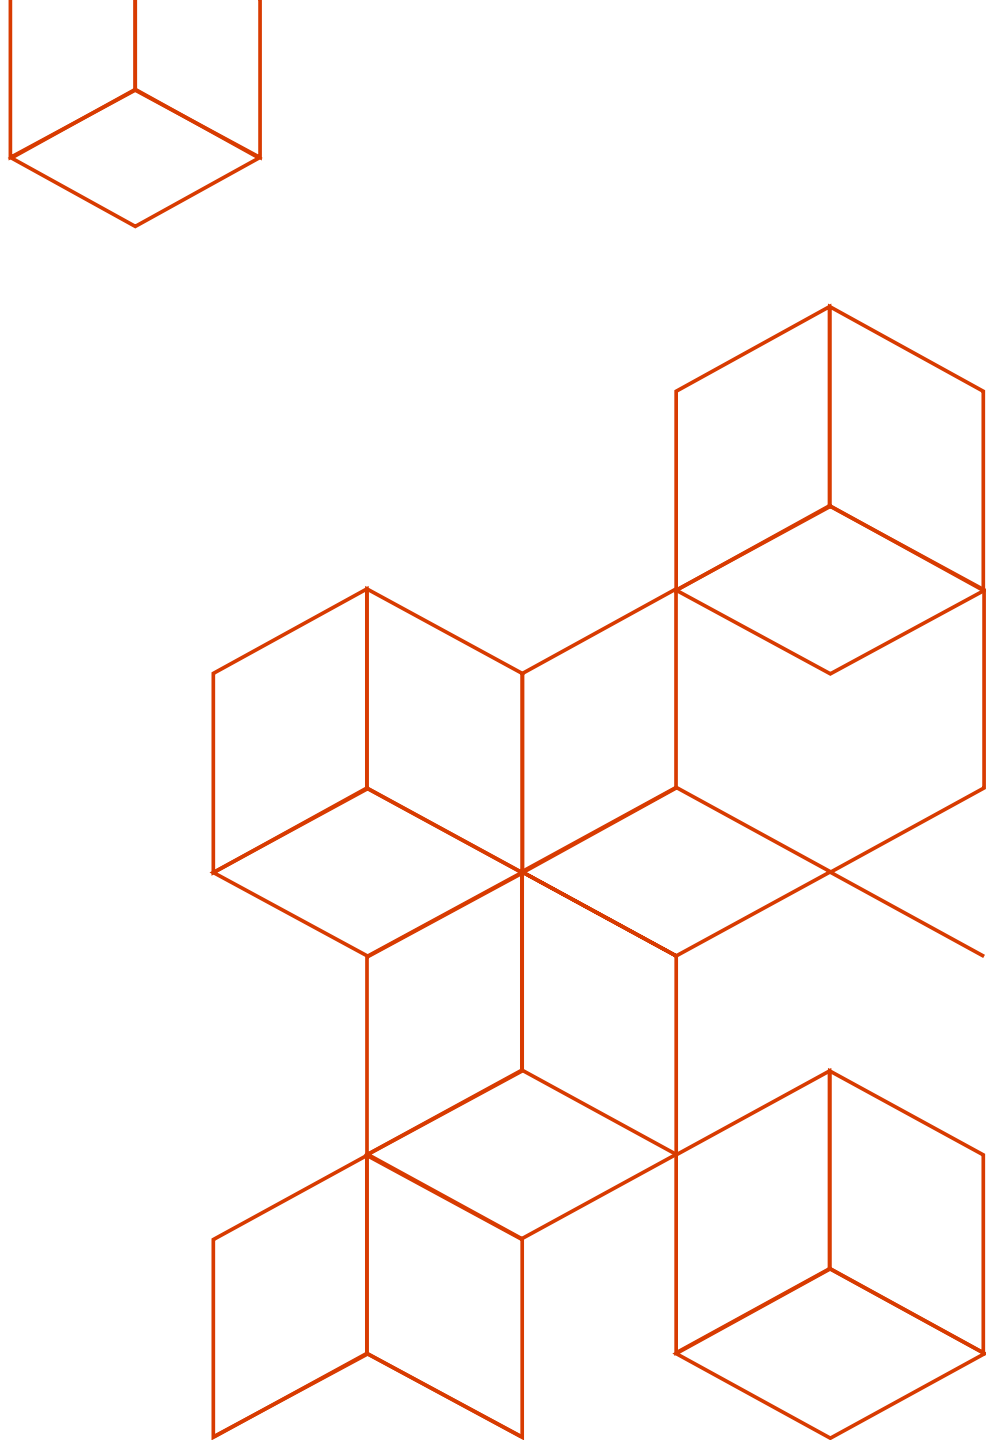

Azure Machine Learning service

Bring AI to everyone with an end-to-end, scalable, trusted platform

Boost your data science productivity

Increase your rate of experimentation

Deploy and manage your models everywhere

Built with your needs in mind

- Automated machine learning
- Managed compute
- Simple deployment
- DevOps for machine learning
- Support for open source frameworks
- Tool agnostic Python SDK

Seamlessly integrated with the Azure Portfolio

Resources beyond this AI Airlift

Azure Machine Learning with Azure Databricks
https://aka.ms/aml-notebook-databricks-e2e

Azure Notebooks
https://notebooks.azure.com/azureml/projects/azureml-getting-started

Azure ML Docs
https://docs.microsoft.com/en-us/azure/machine-learning/service/

Questions

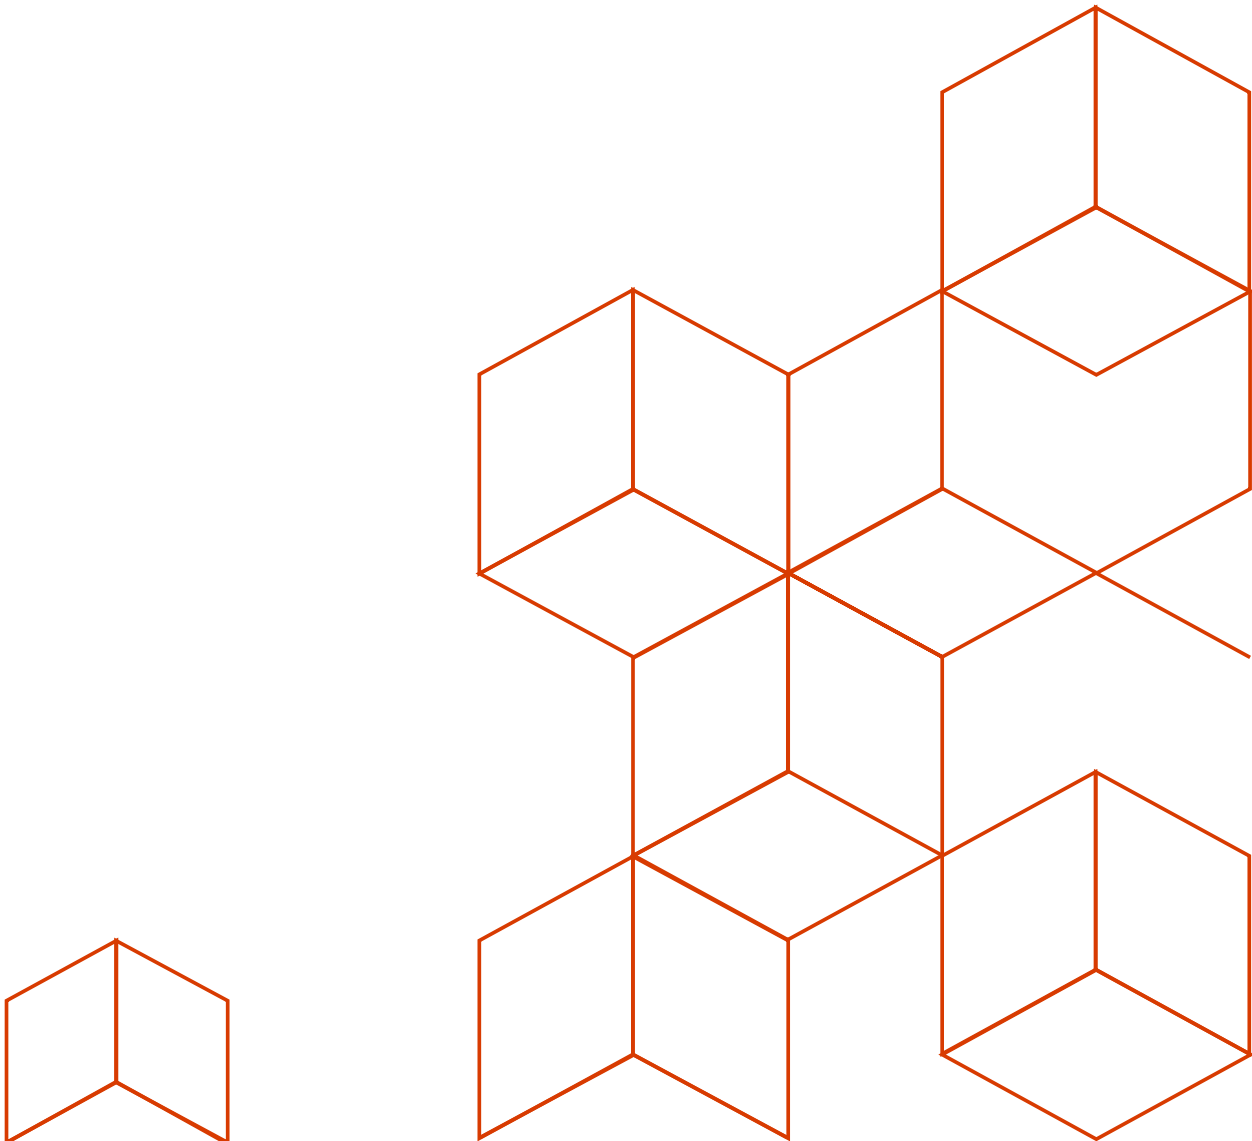

Resources for this Airlift

Azure Subscriptions
https://aka.ms/learnAIsubscriptions

Azure Databricks Notebooks
https://aka.ms/learnAInotebooks.dbc

Git Repository for LearnAI CustomAI Partner Airlift
https://github.com/azure/learnai_azure_ml